




The first spring 2021 assignment explained

Jehan-François Pâris
jfparis@uh.edu

The model

We have

- One single core CPU
- One disk
- One print spooler
- Three queues
 - CPU (Ready queue)
 - Disk
 - Spooler



A deck of punched cards





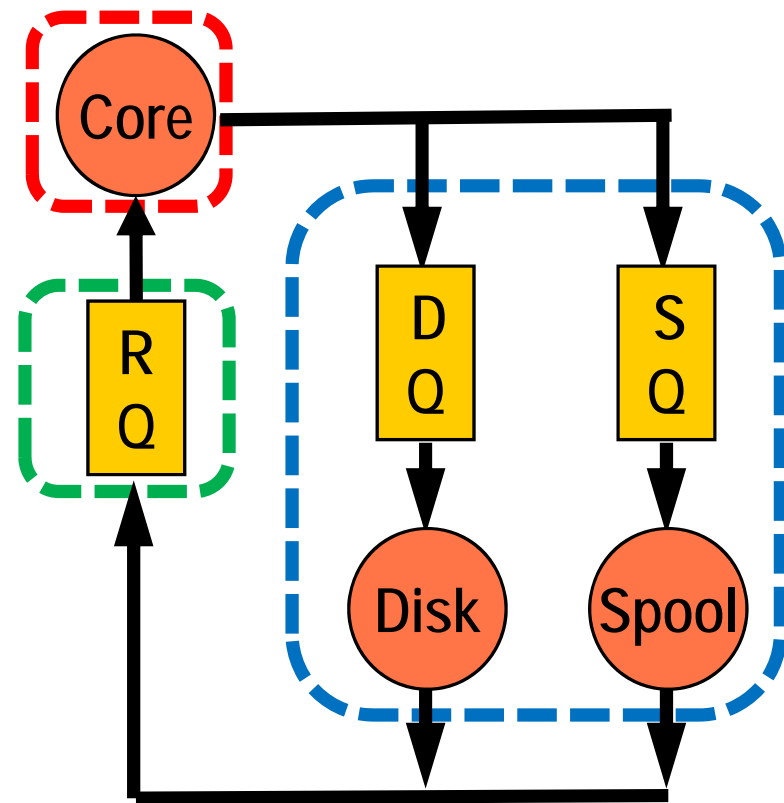
A very simple case

- | | | |
|-------|------|---------------------------|
| MPL | 1 | // no multiprogramming |
| JOB | 1 | // new job |
| CORE | 100 | // request CORE for 100ms |
| DISK | 0 | // no wait disk request |
| CORE | 30 | // request CORE for 30ms |
| DISK | 7 | // request DISK for 7ms |
| CORE | 20 | // request CORE for 20s |
| PRINT | 1000 | // print spooler request |
| CORE | 20 | // request CORE for 30ms |

Job/process states

A job can be

- **Running**
 - ❑ It occupies a core
- **Ready**
 - ❑ It waits for a core
- **Blocked**
 - ❑ It waits for the completion of a system request





Starting the simulation

- **MPL 1**
 - Jobs are processed one by one
 - Set (simulated) time to zero
- **JOB 1**
 - Fetch and start processing job



Output


- Job 1 starts at time 0 ms

Job Table:

There are no other active jobs

Allocate core to job 1 for 100ms at $t = 0\text{ms}$

■ MPL 1
JOB 1
CORE 100
DISK 0
CORE 30
DISK 7
CORE 20
PRINT 1000
CORE 20






What's next?

- Must wait end of core step at $t = 100\text{ms}$
- Can then perform the next step

- **Set $t = 100\text{ms}$**

Zero-time disk access at $t = 100\text{ms}$

■ MPL 1
JOB 1
CORE 100
DISK 0
CORE 30
DISK 7
CORE 20
PRINT 1000
CORE 20





What's next?


- Immediately perform the next step
- At $t = 100\text{ms}$

A zero-delay disk request (DISK 0)

- ☐ ***Always bypasses*** the disk
- ☐ ***Never waits*** for it

Allocate core to job 1 for 30ms at $t = 100\text{ms}$

■ MPL	1
JOB	1
CORE	100
DISK	0
<u>CORE</u>	<u>30</u>
DISK	7
CORE	20
PRINT	1000
CORE	20






What's next?

- Must wait end of core step at $t = 100\text{ms} + 30\text{ms} = 130\text{ms}$
- Can then perform the next step
- **Set $t = 130\text{ms}$**

Allocate disk to job 1 for 7ms at $t = 130\text{ms}$

■ MPL	1
JOB	1
CORE	100
DISK	0
CORE	30
DISK	7
CORE	20
PRINT	1000
CORE	20






What's next?

- Must wait end of disk step at $t = 130\text{ms} + 7\text{ms} = \underline{137\text{ms}}$
- Can then perform the next step

- **Set $t = 137\text{ms}$**

Allocate core to job 1 for 20ms at $t = 137\text{ms}$

■ MPL	1
JOB	1
CORE	100
DISK	0
CORE	30
DISK	7
CORE	20
PRINT	1000
CORE	20






What's next?

- Must wait end of core step at $t = 137\text{ms} + 20\text{ms} = \underline{157\text{ms}}$
- Can then perform the next step

□ **Set $t = 157\text{ms}$**

Allocate spooler to job 1 for 1000ms
at $t = 157\text{ms}$

■ MPL	1
JOB	1
CORE	100
DISK	0
CORE	30
DISK	7
CORE	20
<u>PRINT</u>	<u>1000</u>
CORE	20






What's next?

- Must wait end of spooler step at $t = 157\text{ms} + 1000\text{ms} = \underline{\mathbf{1157\text{ms}}}$
- Can then perform the next step

- **Set $t = 1157\text{ms}$**

Allocate core to job 1 for 20ms at $t = 1157\text{ms}$

■ MPL 1
JOB 1
CORE 100
DISK 0
CORE 30
DISK 7
CORE 20
PRINT 1000
CORE 20





What's next?

- Job 1 has terminated
- Print the required output



Output

- Job 1 terminates at time 1177ms

Job Table:

There are no active jobs

- SUMMARY:

Total elapsed time: 1177ms

Number of jobs that have completed: 1

Total number of disk accesses: 2

Core utilization: 0.13



Things become interesting

```
■ MPL      2          // memory can hold 2 jobs
  JOB      1          // new job
  CORE     100        // request CORE for 100ms
  DISK      9        // request DISK for 9ms
  CORE     30        // request CORE for 30ms
  JOB      5          // new job
  CORE     20        // request CORE for 20s
  DISK      8        // request DISK for 8ms
  CORE     20        // request CORE for 30ms
```



Starting the simulation

- **MPL 2**
 - Two jobs now compete for system resource
 - Better utilization of CPU, disk, ...
 - Price to pay is ***queuing delays***
 - Set (simulated) time to zero
- **JOB 1**
 - Fetch and start processing job 1
- **JOB 5**
 - Fetch and put on hold job 5 in ready queue



Output

- Job 1 starts at time 0 ms

Job Table:

There are no other active jobs


- Job 5 starts at time 0 ms

Job Table:

Job 1 is RUNNING (Will also accept Job 1 is READY)

Allocate core to job 1 for 100ms at $t = 0$ ms
Note job 5 requests core

■ MPL	2
JOB	1
<u>CORE</u>	<u>100</u>
DISK	9
CORE	30
JOB	5
<u>CORE</u>	<u>20</u>
DISK	8
CORE	20






What's next?

- Must wait end of core step at $t = 0\text{ms} + 100\text{ms} = \underline{100\text{ms}}$
- Can then perform the next step

- **Set $t = 100\text{ms}$**

Allocate disk to job 1 for 9ms at $t = 100\text{ms}$
Allocate core to job 5 for 20ms

■ MPL	2
JOB	1
CORE	100
<u>DISK</u>	<u>9</u>
CORE	30
JOB	5
<u>CORE</u>	<u>20</u>
DISK	8
CORE	20





What's next

- Must wait for the first task to complete:
 - Disk I/O for job 1 ending at $t = 100\text{ms} + 9\text{ms} = \underline{\mathbf{109\text{ms}}}$
 - CORE request from job 5 ending at $t = 100\text{ms} + 20\text{ms} = 120\text{ms}$
- Disk I/O for job 1 completes first
 - **Set $t = 109\text{ms}$**



Event list


- Tells your program decide what step to take next
- Each record describes a completion event
 - Expected time
 - Type of event
 - Job ID

$t = 109\text{ms}$
DISK
Job 1

$t = 120\text{ms}$
CORE
Job 5


Using the event list

- When deciding which step to take, must always pick the one associated with the next event



Job 1 requests a core at $t = 109\text{ms}$

■ MPL	2
JOB	1
CORE	100
DISK	9
<u>CORE</u>	<u>30</u>
JOB	5
<u>CORE</u>	<u>20</u>
DISK	8
CORE	20



The new event list


- Only one entry
- **Set $t = 120\text{ms}$**

**$t = 120\text{ms}$
CORE
Job 5**



Allocate core to job 1 for 30ms at $t = 120\text{ms}$
Allocate disk to job 5 for 8ms

■ MPL	2
JOB	1
CORE	100
DISK	9
<u>CORE</u>	<u>30</u>
JOB	5
CORE	20
<u>DISK</u>	<u>8</u>
CORE	20






What's next

- Must wait for the first task to complete:
 - Disk I/O for job 5 ending at $t = 120\text{ms} + 8\text{ms} = \underline{128\text{ms}}$
 - CORE request from job 1 ending at $t = 120\text{ms} + 30\text{ms} = 150\text{ms}$
- Disk I/O for job 5 completes first
 - **Set $t = 128\text{ms}$**


The new event list

- When deciding which step to take, must always pick the one associated with the next event



Job 5 requests CORE at $t = 128\text{ms}$

■ MPL	2
JOB	1
CORE	100
DISK	9
<u>CORE</u>	<u>30</u>
JOB	5
CORE	20
DISK	8
<u>CORE</u>	<u>20</u>






The new event list

$t = 150\text{ms}$
CORE
Job 1



Job 1 terminates at $t = 150\text{ms}$
Allocate core to job 5 for 20ms

■ MPL	2
JOB	1
CORE	100
DISK	9
CORE	30
JOB	5
CORE	20
DISK	8
<u>CORE</u>	<u>20</u>





Output

- Job 1 terminates at time 150ms

Job Table:

Job 5 is RUNNING



Output

- Job 1 terminates at time 150ms

Job Table:

Job 5 is RUNNING

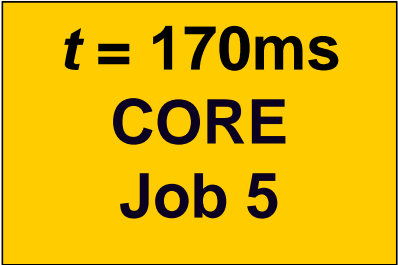


What's next

- Must wait for the completion of core request for job 5 ending at $t = 150\text{ms} + 20\text{ms} = \underline{170\text{ms}}$
- **Set $t = 170\text{ms}$**



The new event list

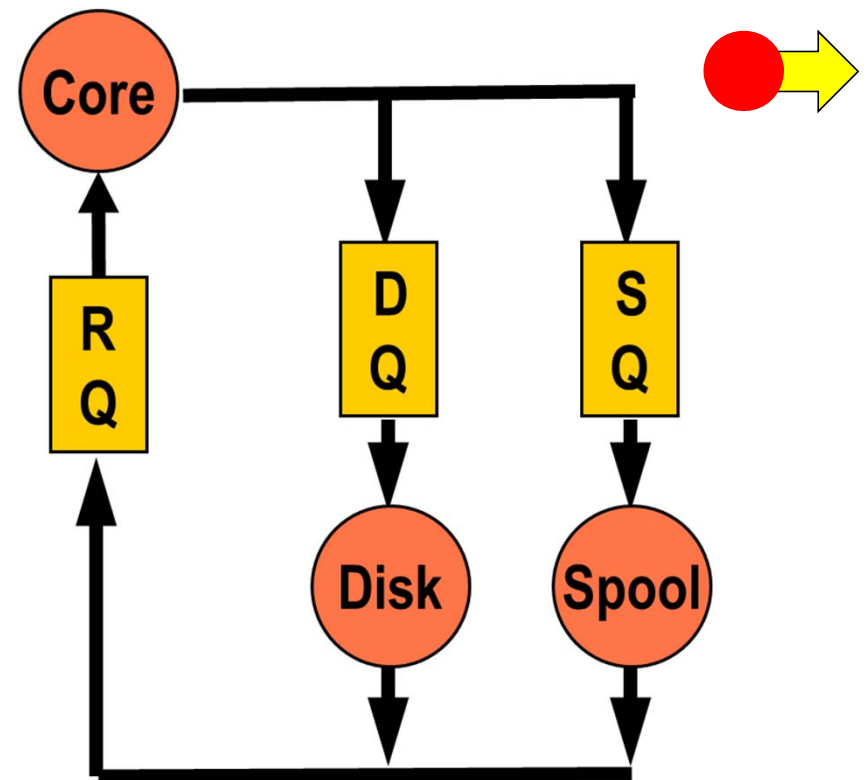


***t* = 170ms**
CORE
Job 5



Job 5 terminates at $t = 170\text{ms}$

■ MPL 2
JOB 1
CORE 100
DISK 9
CORE 30
JOB 5
CORE 20
DISK 8
CORE 20





Output

- Job 5 terminates at time 170ms

Job Table:

There are no active jobs

- SUMMARY:

Total elapsed time: 170ms

Number of jobs that have completed: 2

Number of disk accesses: 2

Core utilization: 1.00



Computing core utilization

- $$\frac{\text{Sum of all CORE times}}{\text{Simulated elapsed time}}$$
- $(100 + 30 + 20 + 20)/170 = 1.00$
 - *The CPU was always busy*



Handling parallel activities

- We only need to consider start times and completion times of each computational step
- Completion times are the most important
 - Release a device
 - Initiate the next request
 - Can be immediately satisfied if requested device is free
 - May require job to wait for device
 - Ready queue, disk queue, spooler queue



The simulated time

- Imaginary clock keeping track of simulated time
 - Never ticks
 - Only updated each time we process a new event



ENGINEERING THE SIMULATION



Simulating time

- ***Absolutely nothing happens to our model between two successive "events"***
- ***Events are***
 - Completion of a computing step
 - Completion of a disk access
 - Completion of a spooler request
- We associate an ***event routine*** with each event

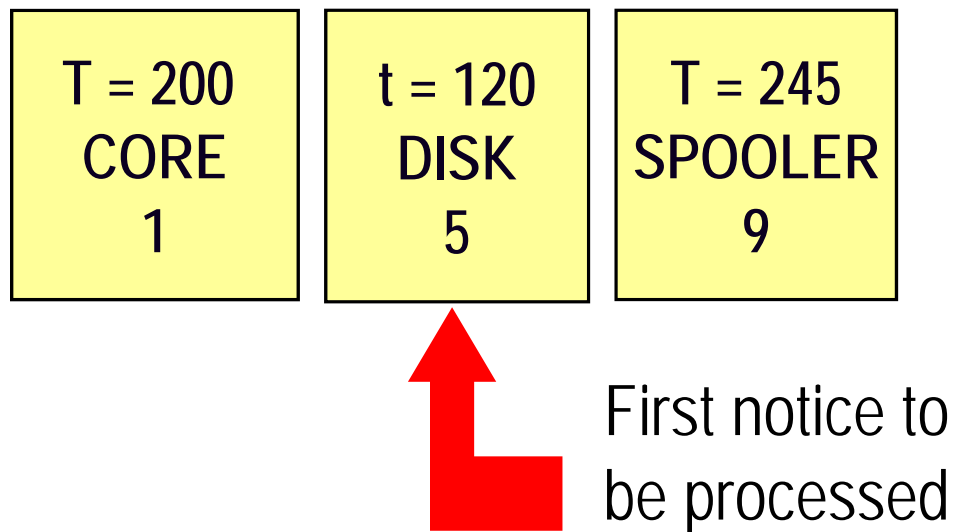


Organizing our program (I)

- Most steps of simulation involve scheduling future completion events
- Associate with each completion event an event notice
 - ***Time of event***
 - ***Device (core, disk, spooler)***
 - ***Job ID***

Organizing our program (II)

- Process all event notices in chronological order





Organizing our program (III)

- Overall organization of main program

```
read in input file
start first MPL jobs
while (event list is not empty) {
    process next event in list
} // while
print simulation results
```



Organizing our program (IV)

- Processing next event in list

```
pop event from list
clock = event.time
if (event.type is core) {
    core(event.time, event.jobID)
else if (event.type is disk) {
    disk(event.time, event.jobID)
```





Organizing our event list (I)

- As a priority queue
- Associating a completion time
 - With each core request
 - With each disk request
 - With each spooler request

Organizing our event list

- Process all event notices in time order



First notice to be processed
is at the head of the list



Core request routine

```
core_request(how_long, jobID){  
    if (core == FREE) {  
        core = BUSY;  
        schedule CORE completion at time  
        current_time + how_long for job jobID;  
    } else {  
        queue jobID in readyQueue  
    } // if  
} // core_request
```



Core completion routine

```
core_release (jobID){  
    if (readyQueue is not empty) {  
        pop first core request in readyQueue  
        schedule its completion at  
        current_time + how_long  
    } else {  
        core = FREE;  
    } //if  
    process next job request for job jobID  
} // core_release
```



Disk request routine

```
disk_request(how_long, jobID){  
    if (how_long == 0) {  
        perform next job request  
    } else if (disk == FREE) {  
        disk = BUSY;  
        schedule DISK completion event at time  
        current_time + how_long for job jobID  
    } else {  
        queue job jobID in diskQueue;  
    } // if  
} // disk_request
```



Disk completion routine

```
disk_release (jobID){  
    if (disk queue is not empty) {  
        pop first request in disk queue  
        schedule its completion at  
        current_time + how_long;  
    } else {  
        disk = FREE;  
    } // if  
    process next job request for job jobID  
} // disk_release
```



Spooler request routine

```
spooler_request(how_long, jobID){  
    if (spooler == FREE) {  
        disk = BUSY;  
        schedule SPOOLER completion event at time  
            current_time + how_long for job jobID  
    } else {  
        queue job jobID in spoolerQueue;  
    } // if  
} // spooler_request
```



Spooler completion routine

```
spooler_release (jobID){  
    if (spooler queue is not empty) {  
        pop first request in spooler queue  
        schedule its completion at  
        current_time + how_long;  
    } else {  
        spooler = FREE;  
    } // if  
    process next job request for job jobID  
} // disk_release
```



Overview (I)

■ ***Input module***

- Read in all input data
 - Store them in a jobList
- Start first MPL jobs

■ ***Main loop***

- Pops next event from event list
 - CORE completion
 - DISK completion
 - SPOOLER completion



Overview (II)

- ***Starting a job***

- Will always be a Core request

- ***Handling job termination***

- ```
if (jobList is empty) {
 break; // then print summary
} else {
 pop next job from jobList
} // if
```





# Overview (III)

- ***Core request***

- If a core is free

- Schedules a CORE completion event

- ***CORE completion event***

- May schedule a CORE completion event

- Starts next request

- DISK or SPOOLER



# Overview (IV)

- ***Disk request***

- If disk is free

- Schedules an DISK completion event

- ***Disk completion event***

- May schedule a DISK completion event

- Starts next request

- Always a CORE request




# Overview (V)

- ***Spooler request***

- Schedules a SPOOLER completion event

- ***Spooler completion event***

- May schedule a SPOOLER completion event
  - Starts next request
    - Always a CORE request





# Explanations

- Green boxes represent conventional functions
- Amber boxes represent events and their associated functions
- Continuous blue arrows represent regular function calls
- Red dashed lines represent the scheduling of specific events



# Finding the next event

- If you do not use a priority list for your events, you can find the next event to process by searching the **lowest value** among the times
  - The CPU will be done with the current job
  - The disk will complete a disk I/O
  - The spooler will be done with the current job



# AN IMPLEMENTATION

- My main data structures would include:
  - An input table
  - A job/process table
  - A device table




# The input table

- Stores the input data
- Line indices are used in job table

| Operation | Parameter |
|-----------|-----------|
| MPL       | 5         |
| JOB       | 3         |
| CORE      | 20        |
| DISK      | 0         |
| CORE      | 20        |
| JOB       | 42        |
| CORE      | 4         |
| ...       | ...       |



# A more elegant input table



Top array indexed  
by job sequence  
numbers, not job ID



## The job table/process table (I)

| Job ID | First Line | Last Line | Current Line | State  |
|--------|------------|-----------|--------------|--------|
| 3      | 1          | 4         | varies       | varies |
| 4      | 5          | ...       | ...          |        |
|        | ...        | ...       | ...          |        |



## The job table/process table (II)

- One line per job
  - ***Line index is job sequence number!***
- First column has start time of job
- First line, last line and current line respectively identify first line, last line and current line of the process in the input table
- Last column is for the current state of the process (READY, RUNNING or BLOCKED)



## The device table (I)

| Device | Status | Busy times total |
|--------|--------|------------------|
| CPU    | P0     | 15               |
| disk   | -      | --               |



## The device table (II)

- One line per device
  - ***Line index identifies the device***
- First column has status of device
  - Number of free cores for CPU
  - Free/busy for disk
- Last column is for the total of all busy times



# Reading your input

- You must use I/O redirection

- `assign1 < input_file`

- ***Advantages***

- Very flexible

- Programmers write their code as if it was reading from standard input

- No need to mess with `fopen()`, `argc` and `argv`



# Detecting the end of data

- The easiest ways to do it
- If you use `scanf()`
  - `scanf(...)` returns `0` once it reaches the end of data
    - `while(scanf(...)) { ... }`
- If you use `cin`
  - `cin` returns `0` once it reaches the end of data
    - `while (cin >> keyword >> argument) { ... }`