



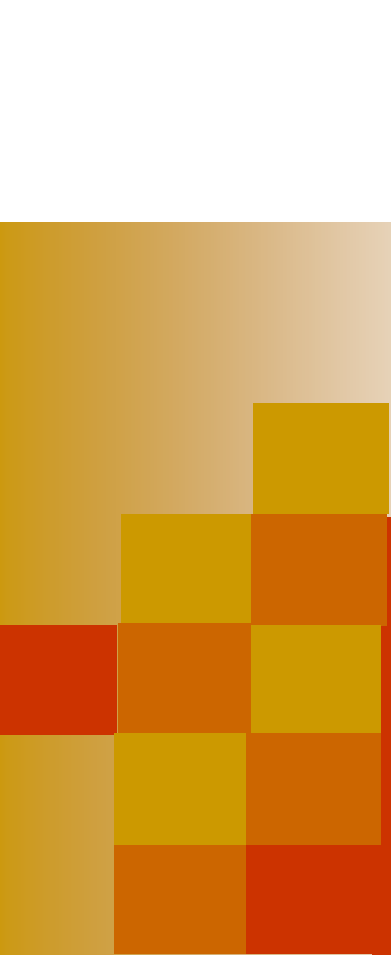
COSC 3360/6310

Monday, January 25



Announcements

- The first assignment is online
 - Folder ***First assignment*** of ***Assignments Channel*** on ***Teams***
 - Sub-folder ***Files*** of ***Resources*** folder on ***Prulu***
- ***Explanations*** will come on Wednesday
- Watch for ***sample inputs*** and ***outputs*** all by next Monday



The four missions (continued)



Functions of an OS

- **Four** basic functions

- ☐ To provide a better user interface
- ☐ To manage the system resources
- ☐ To protect users' programs and data
- ☐ To let programs exchange information



Managing system resources

- ***Focus of the remainder of the course***
- ***Not an easy task***
 - Enormous gap between CPU speeds and disk access times



The memory hierarchy (I)

Level	Device	Access Time
1	Fastest registers (2 GHz)	0.5 ns
2	Main memory	10-70 ns
3	Secondary storage (flash)	35-100 μs
4	Secondary storage (disk)	3-12 ms
5	Mass storage (off line)	a few s



The memory hierarchy (II)

- To make sense of these numbers, let us consider an analogy



Writing a paper (I)

Level	Resource	Access Time
1	Open book on desk	1 s
2	Book on desk	
3	Book in UH library	
4	Book in another library	
5	Book very far away	



Writing a paper (II)

Level	Resource	Access Time
1	Open book on desk	1s
2	Book on desk	20-140 s
3	Book in UH library	
4	Book in another library	
5	Book very far away	



Writing a paper (II)

Level	Resource	Access Time
1	Open book on desk	1s
2	Book on desk	20-140s
3	Book in UH library	20-55h
4	Book in another library	
5	Book very far away	



Writing a paper (III)

Level	Resource	Access Time
1	Open book on desk	1 s
2	Book on desk	20-140 s
3	Book in UH library	20-55 h
4	Book in another library	70-277 days
5	Book very far away	



Writing a paper (V)

Level	Resource	Access Time
1	Open book on desk	1 s
2	Book on desk	20s-140 s
3	Book in UH library	20-55 h
4	Book in another library	70-277 days
5	Book very far away	> 63 years



Will the problem go away?

- New storage technologies
 - Cheaper than main memory
 - Faster than disk drives
- Flash drives
- Optane memory



Flash drives

- Offspring of EEPROM memories
- Fast reads
 - Block-level
- Slower writes
 - Whole page of data must be erased then rewritten
- Can only go through a finite number of program /erase cycles



Optane memory (I)

- Byte-addressable non-volatile memory (BNVM)
- Simpler design
 - Bits are stored as resistivity levels of a secret alloy
 - No transistors (≠ SRAM and DRAM)
- Faster than flash
 - 100-300 ns



Optane memory (II)

- Now

- ☐ Non-volatile RAM
- ☐ Disk cache

- In a few years

- ☐ Could replace flash (phones, laptops, ...)
- ☐ Flash could replace disks (disk farms)
- ☐ Disks could replace slower devices
- ☐ Will require a ***redesign*** of file system



Optimizing disk accesses

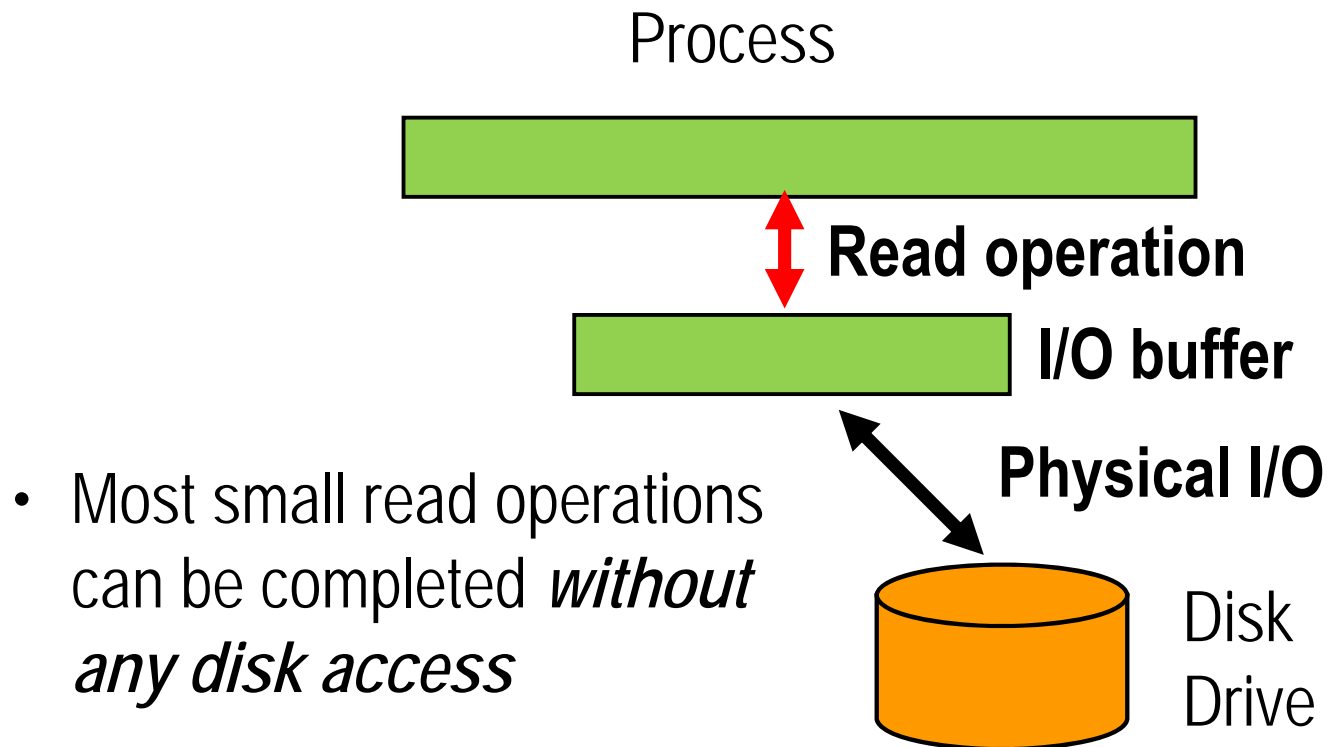
- Two main techniques
 - Making disk accesses more efficient
 - Doing something else while waiting for an I/O operation
- *Not very different from what we are doing in our every day's lives*



Optimizing read accesses (I)

- *When we shop in a market that's far away from our home, we plan ahead and buy food for several days*
- The OS will read as many bytes as it can during each disk access
 - In practice, entire blocks (4KB or more)
 - Blocks are stored in the I/O buffer

Optimizing read accesses (II)





Optimizing read accesses (III)

- Buffered reads work quite well
 - Most systems use it
- Major limitation
 - Cannot read ***too much ahead*** of the program
 - Could end bringing into main memory data that would ***never be used***



Optimizing read accesses (IV)

- Can ***also keep*** in a buffer recently accessed blocks hoping they will be accessed again
 - ***Caching***
- Works very well because we keep accessing again and again the data we are working with
- ***Caching is a fundamental technique of OS and database design***



Optimizing write accesses (I)

- *If we live far away from a library, we wait until we have several books to return before making the trip*
- The OS will **delay writes** for a few seconds then write an entire block
 - Since most writes are sequential, most small writes will not require any disk access



Optimizing write accesses (II)

- ***Delayed writes*** work quite well
 - Most systems use it
- ***Major drawback***
 - We will ***lose data*** if the system or the program crashes
 - After the program issued a write but
 - Before the data were saved to disk
 - Unless we use NVRAM



Doing something else

- *When we order something on the web, we do not remain idle until the goods are delivered*
- The OS can implement ***multiprogramming*** and let the CPU run another program while a program waits for an I/O



Advantages (I)

- Multiprogramming is very important in business applications
 - Many of these applications use the peripherals much more than the CPU
 - For a long time the CPU was the most expensive component of a computer
 - ***Multiprogramming*** was invented to keep the CPU busy



Advantages (II)

- Multiprogramming made ***time-sharing*** possible
- Multiprogramming lets your PC run several applications at the same time
 - MS Word and MS Outlook



Multiprogramming (I)

- Multiprogramming lets the CPU divide its time among different tasks:
 - One tenth of a second on a program, then another tenth of a second on another one and so forth
- Each core of your CPU will still be working on ***one single task*** at any given time



Multiprogramming (II)

- The CPU does not waste any time waiting for the completion of I/O operations
- From time to time, the OS will need to regain control of the CPU
 - Because a task has exhausted its fair share of the CPU time
 - Because something else needs to be done.
- This is done through ***interrupts***.



Interrupts (I)

- Request to interrupt the flow of execution the CPU
- Detected by the CPU hardware
 - **After** it has executed the current instruction
 - **Before** it starts the next instruction.



A very schematic view (I)

- A very basic CPU would execute the following loop:

```
forever {  
    fetch_instruction();  
    decode_instruction();  
    execute_instruction();  
}
```

- Pipelining makes things more complicated
 - And CPU much faster!



A very schematic view (II)

- We add an extra step:

```
forever {  
    check_for_interrupts();  
    fetch_instruction();  
    decode_instruction();  
    execute_instruction();  
}
```





Interrupts (II)

- When an interrupt occurs:
 - a. The ***current state of the CPU*** (program counter, program status word, contents of registers, and so forth) is saved, normally on the top of a stack
 - b. A ***new CPU state*** is fetched



Interrupts (III)

- New state includes a new ***hardware-defined*** value for the program counter
 - Cannot “hijack” interrupts
- Process is totally transparent to the task being interrupted
 - A process ***never*** knows whether it has been interrupted or not



Types of interrupts (I)

- ***I/O completion interrupts***

- Notify the OS that an I/O operation has completed,

- ***Timer interrupts***

- Notify the OS that a task has exceeded its quantum of core time



Types of interrupts (II)

■ ***Traps***

- Notify the OS of a ***program error*** (division by zero, illegal op code, illegal operand address, ...) or a *hardware failure*

■ ***System calls***

- Notify OS that the running task wants to submit a request to the OS



A surprising discovery

- ***Programs do interrupt themselves!***



Context switches

- Each interrupt will result into ***two context switches***:
 - One when the running task is interrupted
 - Another when it regains the CPU
- Context switches are ***not cheap***
- The overhead of any simple system call is ***two context switches***

Remember that!



Prioritizing interrupts (I)

- Interrupt requests may occur while the system is processing another interrupt
- All interrupts are not equally urgent (as it is also in real life)
 - Some are more urgent than other
 - *Also true in real life*



Prioritizing interrupts (II)

- The best solution is to *prioritize* interrupts and assign to each source of interrupts a ***priority level***
 - New interrupt requests will be allowed to interrupt lower-priority interrupts but will have to wait for the completion of all other interrupts
- Solution is known as ***vectorized interrupts***.



Example from real life

- Let us try to prioritize
 - Phone is ringing
 - Washer signals end of cycle
 - Dark smoke is coming out of the kitchen
 - ...
- With vectorized interrupts, a phone call will never interrupt another phone call



The solution

Smoke in the kitchen
Phone is ringing
End of washer cycle
More low-priority stuff



Disabling Interrupts

- We can ***disable*** interrupts
- OS does it before performing short critical tasks that cannot be interrupted
 - Works only for single-threaded kernels
- User tasks ***must*** be prevented from doing it
 - Too dangerous



DMA

- Disk I/O poses a special problem
 - CPU will have to transfer large quantities of data between the disk controller's buffer and the main memory
- ***Direct memory access (DMA)*** allows the disk controller to read data from and write data to main memory without any CPU intervention
 - Controller “steals” memory cycles from CPU