



Chapter I

Introduction

Jehan-François Pâris
jfparis@uh.edu

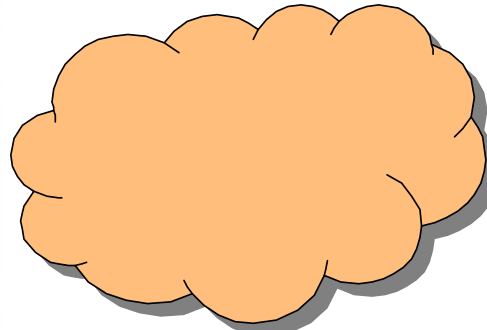


Chapter Overview

- Defining operating systems
- Major functions of an OS
- Types of operating systems
- UNIX
- Kernel organization

What is an operating system?

- “What stands between the user and the bare machine”





What is an operating system?

- The *basic* software required to operate a computer.
- Similar role to that of the conductor of an orchestra



Do not belong to OS

- All ***user programs***
 - Compilers, spreadsheets, word processors, and so forth
 - Most utility programs
 - ***mkdir*** is a user program calling ***mkdir()***
- The ***command language interpreter***
 - Anyone can write his/her UNIX shell



The UNIX shells

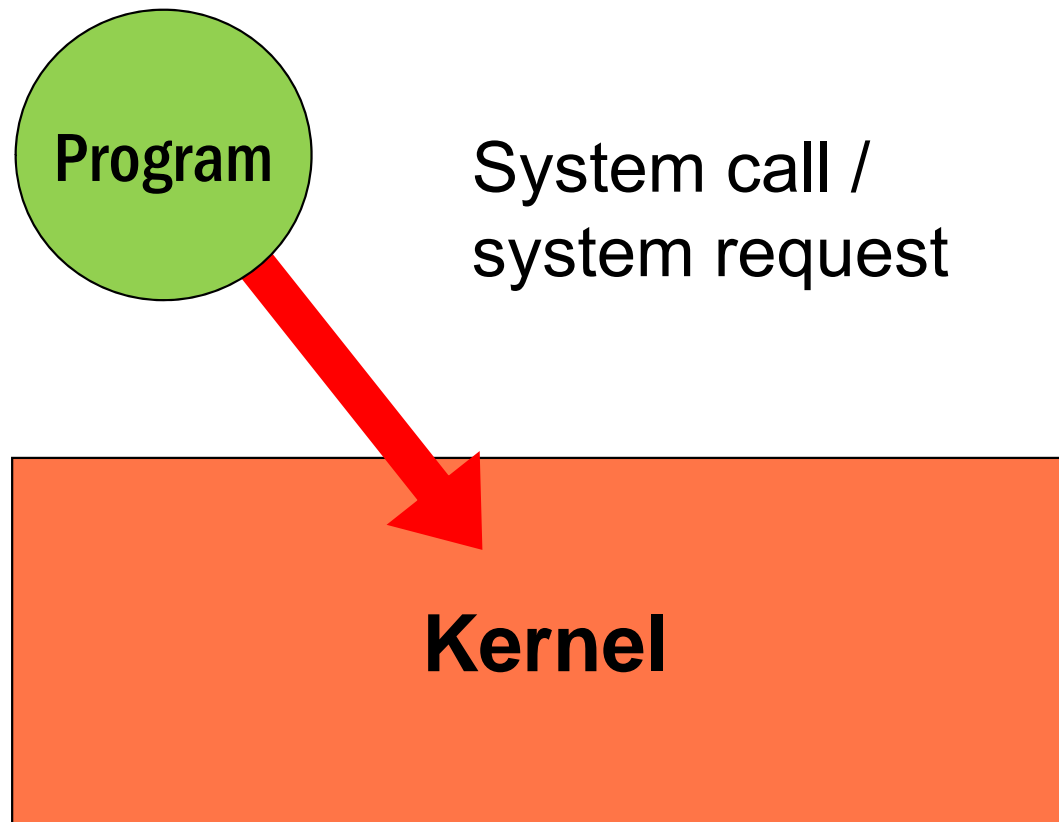
- UNIX has several shells
 - **sh** (the Bourne shell) is the original UNIX shell
 - **cs****h** was developed at Berkeley by Bill Joy
 - **ksh** (the Korn shell) was developed by David Korn at AT&T Bell Laboratories
 - **bash** (the GNU Bourne-Again shell)
- and the list is far from complete



The core of the OS

- Part that remains in main memory
- Controls the execution of all other programs.
- Known as the ***kernel***
 - Also called ***monitor, supervisor, executive***
- Other programs interact with it through **system calls**

System calls





A question

- Who among you has already used system calls?



The answer

- All of you

- ☐ *All I/O operations are performed through system calls*



The four missions



Functions of an OS

- **Four** basic functions

- ☐ To provide a better user interface
- ☐ To manage the system resources
- ☐ To protect users' programs and data
- ☐ To let programs exchange information



A better user interface

- Accessing directly the hardware would be very cumbersome
- Must enter manually the code required to read into main memory each program
 - ***boot strapping***

How it was done (I)



PDP 8

- Early 70's
- 12-bit machine
 - 4K RAM!

How it was done (II)



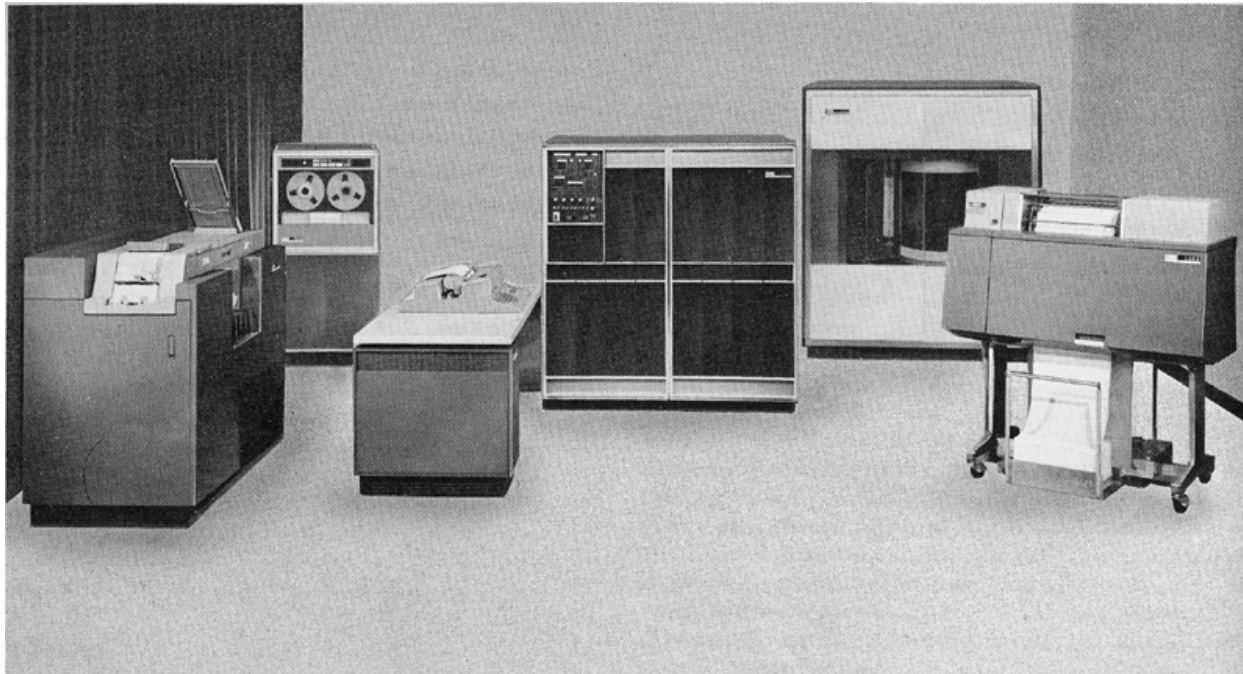
Toggle switches in front panel were used to enter the bootstrap code



Batch systems

- Allow users to submit a batches of requests to be processed in sequence
- Include a command language specifying what to do with the inputs
 - Compile
 - Link edit
 - Execute and so forth

An IBM 1401





Interactive systems

- Came later
- Allow users to interact with the OS through their terminals:
- Include an ***interactive*** command language
 - UNIX shells, Windows PowerShell
 - Can also be used to write scripts



Time sharing

- Lets several interactive users to access a single computer at the same time
- Standard solution when computers were expensive



Graphical user interfaces

- Called GUIs (pronounced *goo-eyes*):
Macintosh, Windows, X-Windows, Linux
 - Require a dedicated computer for each user
 - Pioneered at Xerox Palo Alto Research Center (Xerox PARC)
 - Popularized by the Macintosh
 - Dominated the market with MS Windows

The Xerox Alto





Xerox PARC (I)

- Founded by XEROX in 1970
- Invented
 - Laser printing
 - Ethernet
 - The GUI paradigm
 - Object-oriented programming (Smalltalk)



Xerox PARC (II)

- All their inventions were brought to market by other concerns
- Popular belief is that Xerox management blew it
- In reality
 - Alto workstations were very expensive
 - Smalltalk was very slow
 - Group was too small to deliver a full system



Smart phones and tablets

- Convergence of four trends
 - Cheaper LCD displays
 - Solid-State Storage (SSD)
 - Faster wireless communications
 - Ubiquitous wireless



History repeats itself

- First successful devices introduced by Apple
 - iPod, iPhone, iPad, ...
 - First iPad was underpowered
- Competition soon grows
 - Cheaper Android devices



With a difference!

- Apple did not "steal" the concept from anyone
- iPods, iPhones, iPads were an instant success
 - Reasonably priced



Two models

■ **Apple:**

- Closed ecosystem (***walled garden***)
- Strict controls on app market
- Missing features
 - No file system

■ **Android:**

- Just the opposite
- Lax controls on app market
- Can access the Linux/Android shell

Is this paradise?





Summary

■ ***Six major steps***

- ☐ Bare bone machine
- ☐ Batch systems
- ☐ Timesharing
- ☐ Personal computer
- ☐ Personal computer with GUI
- ☐ Smart phone/tablet



File systems

- Let users create and delete files without having to worry about disk allocation
 - Users lose the ability to specify how their files are stored on the disk
 - Database designers prefer to bypass the file system
- Some file systems tolerate disk failures (RAID)



Managing system resources

- ***Focus of the remainder of the course***
- ***Not an easy task***
 - Enormous gap between CPU speeds and disk access times



The memory hierarchy (I)

Level	Device	Access Time
1	Fastest registers (2 GHz)	0.5 ns
2	Main memory	10-70 ns
3	Secondary storage (flash)	35-100 μs
4	Secondary storage (disk)	3-12 ms
5	Mass storage (off line)	a few s



The memory hierarchy (II)

- To make sense of these numbers, let us consider an analogy



Writing a paper (I)

Level	Resource	Access Time
1	Open book on desk	1 s
2	Book on desk	
3	Book in UH library	
4	Book in another library	
5	Book very far away	



Writing a paper (II)

Level	Resource	Access Time
1	Open book on desk	1s
2	Book on desk	20-140 s
3	Book in UH library	
4	Book in another library	
5	Book very far away	



Writing a paper (II)

Level	Resource	Access Time
1	Open book on desk	1s
2	Book on desk	20-140s
3	Book in UH library	20-55h
4	Book in another library	
5	Book very far away	



Writing a paper (III)

Level	Resource	Access Time
1	Open book on desk	1 s
2	Book on desk	20-140 s
3	Book in UH library	20-55 h
4	Book in another library	70-277 days
5	Book very far away	



Writing a paper (V)

Level	Resource	Access Time
1	Open book on desk	1 s
2	Book on desk	20s-140 s
3	Book in UH library	20-55 h
4	Book in another library	70-277 days
5	Book very far away	> 63 years



Will the problem go away?

- New storage technologies
 - Cheaper than main memory
 - Faster than disk drives
- Flash drives
- Optane memory



Flash drives

- Offspring of EEPROM memories
- Fast reads
 - Block-level
- Slower writes
 - Whole page of data must be erased then rewritten
- Can only go through a finite number of program /erase cycles



Optane memory (I)

- Byte-addressable non-volatile memory (BNVM)
- Simpler design
 - Bits are stored as resistivity levels of a secret alloy
 - No transistors (≠ SRAM and DRAM)
- Faster than flash
 - 100-300 ns



Optane memory (II)

- Now

- ☐ Non-volatile RAM
- ☐ Disk cache

- In a few years

- ☐ Could replace flash (phones, laptops, ...)
- ☐ Flash could replace disks (disk farms)
- ☐ Disks could replace slower devices
- ☐ Will require a ***redesign*** of file system



Optimizing disk accesses

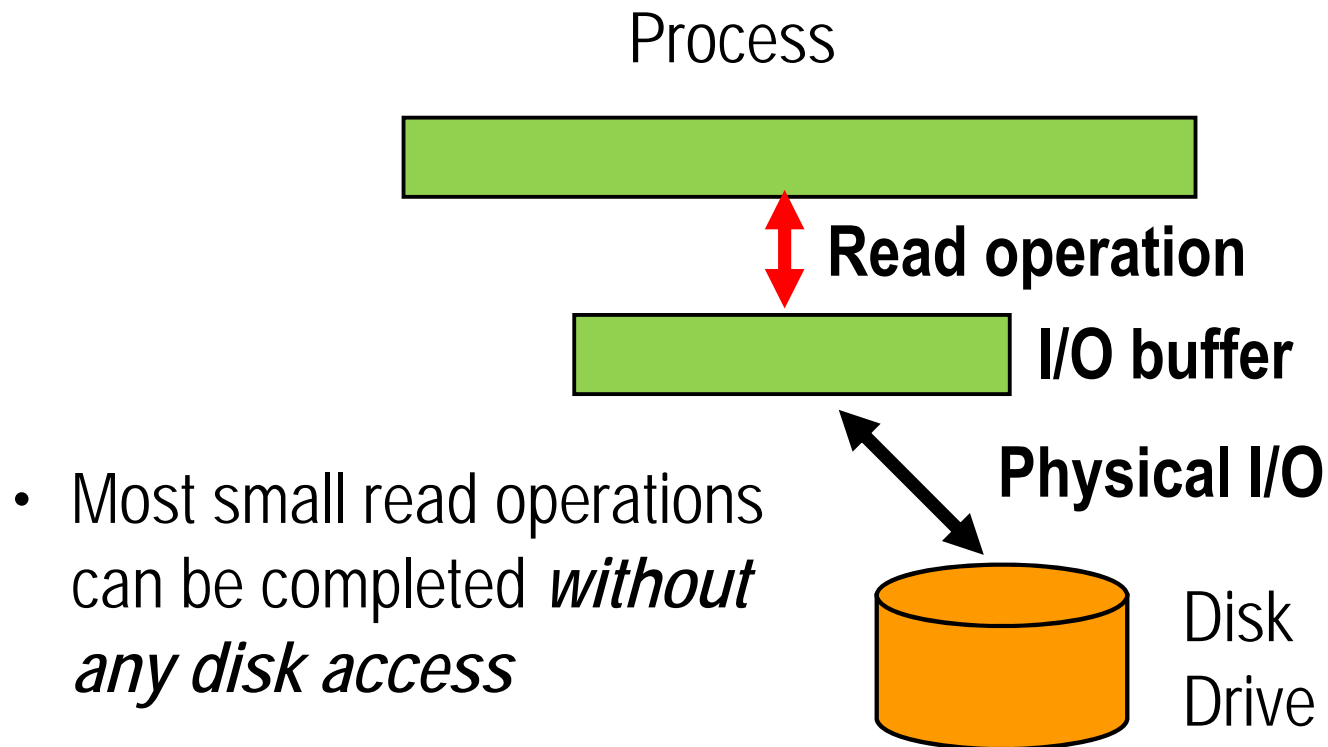
- Two main techniques
 - Making disk accesses more efficient
 - Doing something else while waiting for an I/O operation
- *Not very different from what we are doing in our every day's lives*



Optimizing read accesses (I)

- *When we shop in a market that's far away from our home, we plan ahead and buy food for several days*
- The OS will read as many bytes as it can during each disk access
 - In practice, entire blocks (4KB or more)
 - Blocks are stored in the I/O buffer

Optimizing read accesses (II)





Optimizing read accesses (III)

- Buffered reads work quite well
 - Most systems use it
- Major limitation
 - Cannot read ***too much ahead*** of the program
 - Could end bringing into main memory data that would ***never be used***



Optimizing read accesses (IV)

- Can ***also keep*** in a buffer recently accessed blocks hoping they will be accessed again
 - ***Caching***
- Works very well because we keep accessing again and again the data we are working with
- ***Caching is a fundamental technique of OS and database design***



Optimizing write accesses (I)

- *If we live far away from a library, we wait until we have several books to return before making the trip*
- The OS will **delay writes** for a few seconds then write an entire block
 - Since most writes are sequential, most small writes will not require any disk access



Optimizing write accesses (II)

- ***Delayed writes*** work quite well
 - Most systems use it
- ***Major drawback***
 - We will ***lose data*** if the system or the program crashes
 - After the program issued a write but
 - Before the data were saved to disk
 - Unless we use NVRAM



Doing something else

- *When we order something on the web, we do not remain idle until the goods are delivered*
- The OS can implement ***multiprogramming*** and let the CPU run another program while a program waits for an I/O



Advantages (I)

- Multiprogramming is very important in business applications
 - Many of these applications use the peripherals much more than the CPU
 - For a long time the CPU was the most expensive component of a computer
 - ***Multiprogramming*** was invented to keep the CPU busy



Advantages (II)

- Multiprogramming made ***time-sharing*** possible
- Multiprogramming lets your PC run several applications at the same time
 - MS Word and MS Outlook



Multiprogramming (I)

- Multiprogramming lets the CPU divide its time among different tasks:
 - One tenth of a second on a program, then another tenth of a second on another one and so forth
- Each core of your CPU will still be working on ***one single task*** at any given time



Multiprogramming (II)

- The CPU does not waste any time waiting for the completion of I/O operations
- From time to time, the OS will need to regain control of the CPU
 - Because a task has exhausted its fair share of the CPU time
 - Because something else needs to be done.
- This is done through ***interrupts***.



Interrupts (I)

- Request to interrupt the flow of execution the CPU
- Detected by the CPU hardware
 - **After** it has executed the current instruction
 - **Before** it starts the next instruction.



A very schematic view (I)

- A very basic CPU would execute the following loop:

```
forever {  
    fetch_instruction();  
    decode_instruction();  
    execute_instruction();  
}
```

- Pipelining makes things more complicated
 - And CPU much faster!



A very schematic view (II)

- We add an extra step:

```
forever {  
    check_for_interrupts();  
    fetch_instruction();  
    decode_instruction();  
    execute_instruction();  
}
```





Interrupts (II)

- When an interrupt occurs:
 - a. The ***current state of the CPU*** (program counter, program status word, contents of registers, and so forth) is saved, normally on the top of a stack
 - b. A ***new CPU state*** is fetched



Interrupts (III)

- New state includes a new ***hardware-defined*** value for the program counter
 - Cannot “hijack” interrupts
- Process is totally transparent to the task being interrupted
 - A process ***never*** knows whether it has been interrupted or not



Types of interrupts (I)

- ***I/O completion interrupts***

- Notify the OS that an I/O operation has completed,

- ***Timer interrupts***

- Notify the OS that a task has exceeded its quantum of core time



Types of interrupts (II)

■ ***Traps***

- Notify the OS of a ***program error*** (division by zero, illegal op code, illegal operand address, ...) or a *hardware failure*

■ ***System calls***

- Notify OS that the running task wants to submit a request to the OS



A surprising discovery

- ***Programs do interrupt themselves!***



Context switches

- Each interrupt will result into **two context switches**:
 - One when the running task is interrupted
 - Another when it regains the CPU
- Context switches are **not cheap**
- The overhead of any simple system call is **two context switches**

Remember that!



Prioritizing interrupts (I)

- Interrupt requests may occur while the system is processing another interrupt
- All interrupts are not equally urgent (as it is also in real life)
 - Some are more urgent than other
 - *Also true in real life*



Prioritizing interrupts (II)

- The best solution is to *prioritize* interrupts and assign to each source of interrupts a ***priority level***
 - New interrupt requests will be allowed to interrupt lower-priority interrupts but will have to wait for the completion of all other interrupts
- Solution is known as ***vectorized interrupts***.



Example from real life

- Let us try to prioritize
 - Phone is ringing
 - Washer signals end of cycle
 - Dark smoke is coming out of the kitchen
 - ...
- With vectorized interrupts, a phone call will never interrupt another phone call



The solution

Smoke in the kitchen
Phone is ringing
End of washer cycle
More low-priority stuff



Disabling Interrupts

- We can ***disable*** interrupts
- OS does it before performing short critical tasks that cannot be interrupted
 - Works only for single-threaded kernels
- User tasks ***must*** be prevented from doing it
 - Too dangerous



DMA

- Disk I/O poses a special problem
 - CPU will have to transfer large quantities of data between the disk controller's buffer and the main memory
- ***Direct memory access (DMA)*** allows the disk controller to read data from and write data to main memory without any CPU intervention
 - Controller “steals” memory cycles from CPU



Protecting users' data (I)

- Unless we have an isolated single-user system, we must prevent users from
 - Accessing
 - Deleting
 - Modifyingwithout authorization other people's programs and data



Protecting users' data (III)

- Two aspects
 - Protecting user's files on disk
 - Preventing programs from interfering with each other
- Two solutions
 - Dual-mode CPUs
 - Memory protection



Historical Considerations

- Earlier operating systems for personal computers did not have any protection
 - They were single-user machines
 - They typically ran one program at a time
- Windows 2000, Windows XP, Vista and MacOS X are protected



Protecting users' files

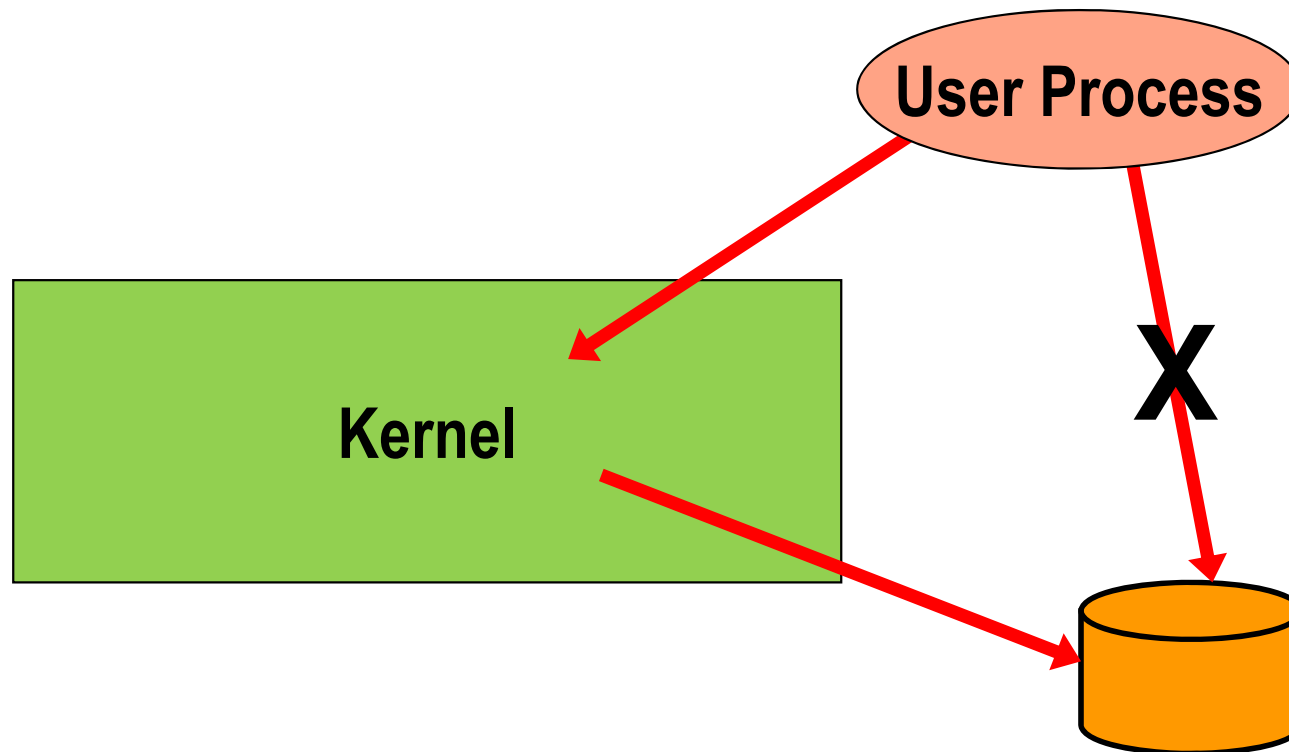
- Key idea is to prevent users' programs from directly accessing the disk
- Will require I/O operations to be performed by the kernel
- Make them ***privileged instructions***
 - Only the kernel can execute



Privileged instructions

- Require a ***dual-mode CPU***
- Two CPU modes
 - ***Privileged mode*** or ***executive mode***
 - Allows CPU to execute all instructions
 - ***User mode***
 - Allows CPU to execute only safe unprivileged instructions
- State of CPU is determined by a ***special bit***

All disk/SSD accesses must go through the kernel

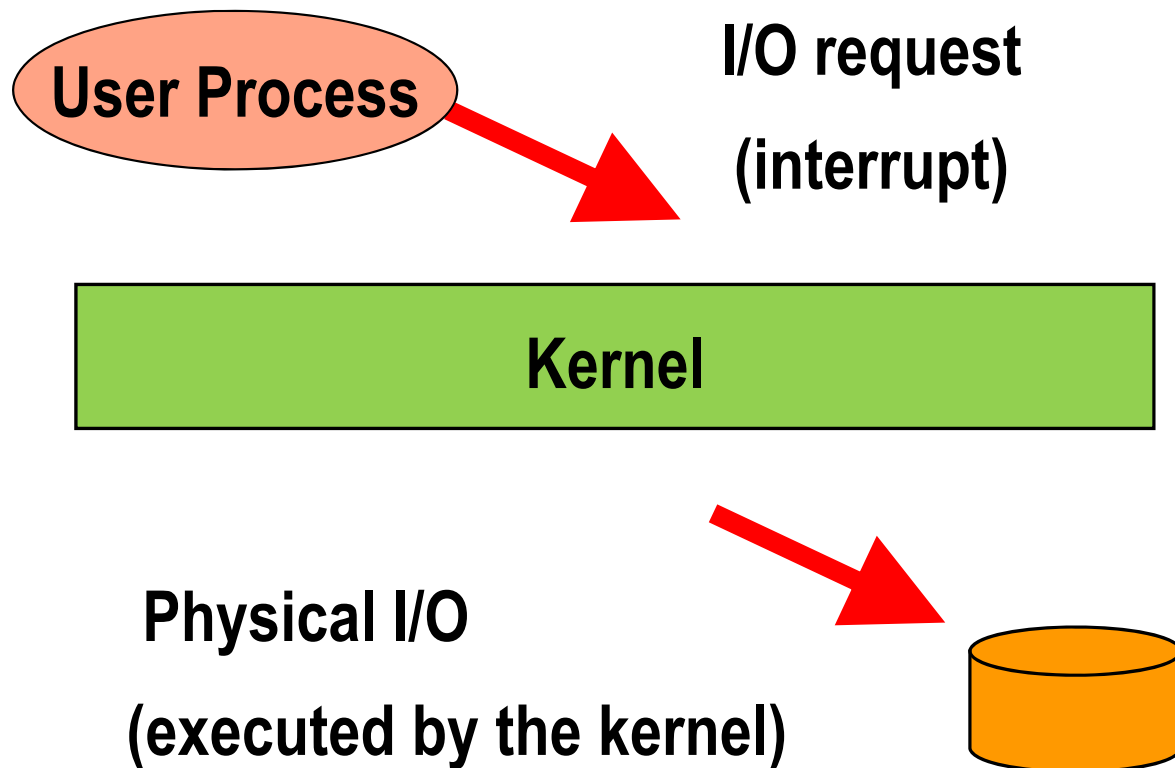




Switching between states

- User mode will be the default mode for all programs
 - Only the kernel can run in supervisor mode
- Switching from user mode to supervisor mode is done through an interrupt
 - Safe because the jump address is at a well-defined location in main memory

Performing an I/O





An analogy (I)

- Most UH libraries are ***open stacks***
 - Anyone can consult books in the stacks and bring them to checkout
- National libraries and the Library of Congress have ***close stack collections***
 - Users fill a request for a specific document
 - A librarian will bring the document to the circulation desk



An analogy (II)

- ***Open stack collections***

- ☐ Let users browse the collections
- ☐ Users can misplace or vandalize books

- ***Close stack collections***

- ☐ Much slower access
- ☐ Much safer



More trouble

- Having a dual-mode CPU is ***not enough*** to protect user's files
- Must also prevent rogue users from tampering with the kernel
 - Same as a rogue customer bribing a librarian in order to steal books
- Done through ***memory protection***



Memory protection (I)

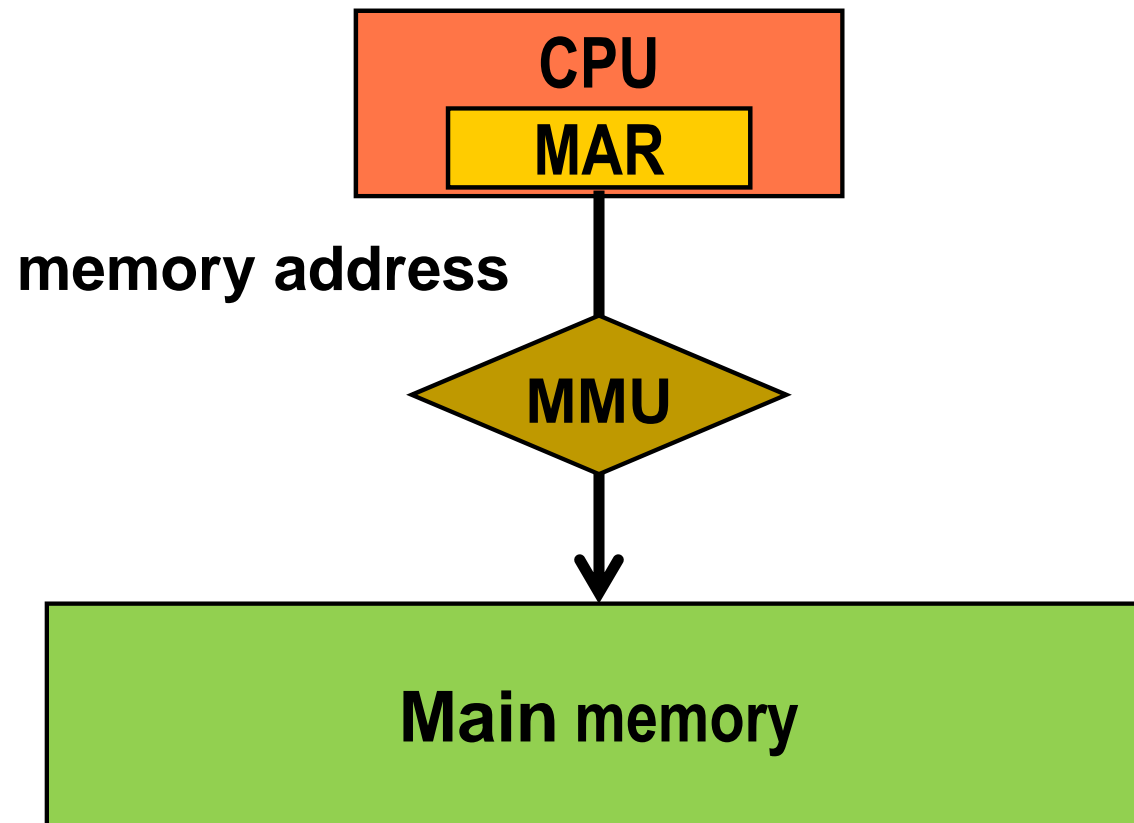
- Prevents programs from accessing any memory location outside their own ***address space***
- Requires special ***memory protection hardware***
 - ***Memory Management Unit (MMU)***
- Memory protection hardware
 - Checks ***every*** reference issued by program
 - Generates an interrupt when it detects a protection violation



Memory protection (II)

- Has additional advantages:
 - Prevents programs from corrupting address spaces of other programs
 - Prevents programs from crashing the kernel
 - Not true for device drivers which are inside the kernel
- Required part of any multiprogramming system

Memory protection (III)





Even more trouble

- Having both a dual-mode CPU and memory protection is not enough to protect user's files
- Must also prevent rogue users from booting the system with a ***doctored kernel***
 - ***Example:***
 - Can run Linux from a “live” CD Linux
 - Linux will read all NTFS files ignoring all restrictions set up by Windows



INTERPROCESS COMMUNICATION

- Has become very important over the last thirty years
- Two techniques
 - Message passing
 - General but not very easy to use
 - Shared memory
 - Less general, easier to use but requires interprocess synchronization



ANOTHER VIEW

- Arpaci-Dusseau & Arpaci-Dusseau
 - Focus on services provided by OSes
- Three themes
 - Virtualization
 - Concurrency
 - Persistence



Virtualization

- The process abstraction
- Virtualizing the CPU:
 - Process scheduling
- Virtualizing the memory:
 - Memory management



Concurrency

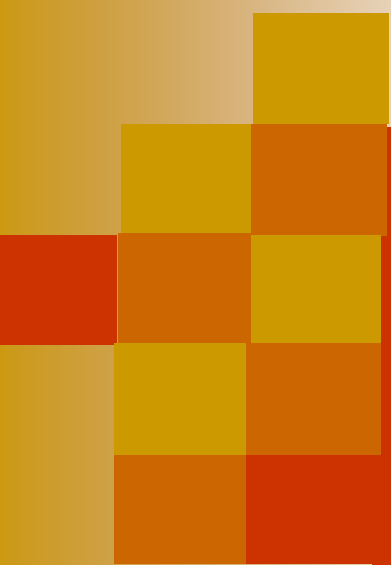
- Threads
- Locks
- Semaphores

We will cover threads in the chapter on processes because they are essential to the client server model



Persistence

- The file system



Types of operating systems



Types of operating systems

- Already discussed:
 - Batch systems
 - Time-Sharing systems
- Will now introduce
 - Real-Time systems
 - Operating systems for multiprocessors
 - Distributed systems



Real-time systems

- Designed for applications with ***strict real-time constraints*** :
 - ***Process control***
 - ***Guidance systems***
 - Most ***multimedia applications***
- Must guarantee that critical tasks will ***always*** be performed within a specific time frame.



Hard RT systems

- Must guarantee that all deadlines will always be met
- ***Any failure*** could have ***catastrophic consequences***:
 - The reactor could overheat and explode
 - The rocket could be lost



Soft RT systems

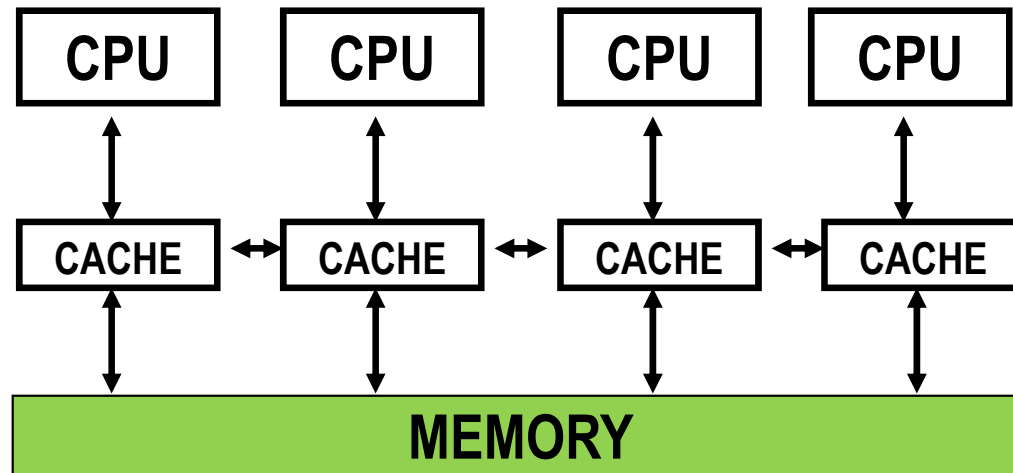
- Guarantee that most deadlines will be met
- A DVD decoder that miss a deadline will spoil our viewing pleasure for a fraction of a second



Observations

- Hard RT applications normally run on special RT OSes
- Soft RT applications can run on a regular OS
 - If the OS supports them
- Interactive and time-sharing systems are ***not*** RT systems
 - They attempt to provide a fast response time but do not try to meet specific deadlines

Multiprocessor operating systems



- Designed for ***multiprocessor architectures***
 - Several processors share the same memory



Leader/follower multiprocessing

- Single copy of OS runs on a dedicated core/processor
 - **Leader** or *master* (deprecated)
- Other cores/processors can only run applications
 - **Followers** or *slaves* (deprecated)
- Major advantage is ***simplicity***
 - Requires few changes
- Major disadvantage is ***lack of scalability***
 - Single copy of OS can become a ***bottleneck***



Symmetric multiprocessing

- Any core/processor can perform all functions
 - There can be multiple copies of the OS running in parallel
- Must prevent them from interfering with each other
 - Disabling interrupts will not work
 - Must add **locks** to all critical sections



The state of the art

- Most computers now have ***multicore CPUs***
 - Sole practical way to increase CPU power
- Many have powerful GPUs
 - Highly parallel
- Using multicore architectures in an effective way is a huge challenge

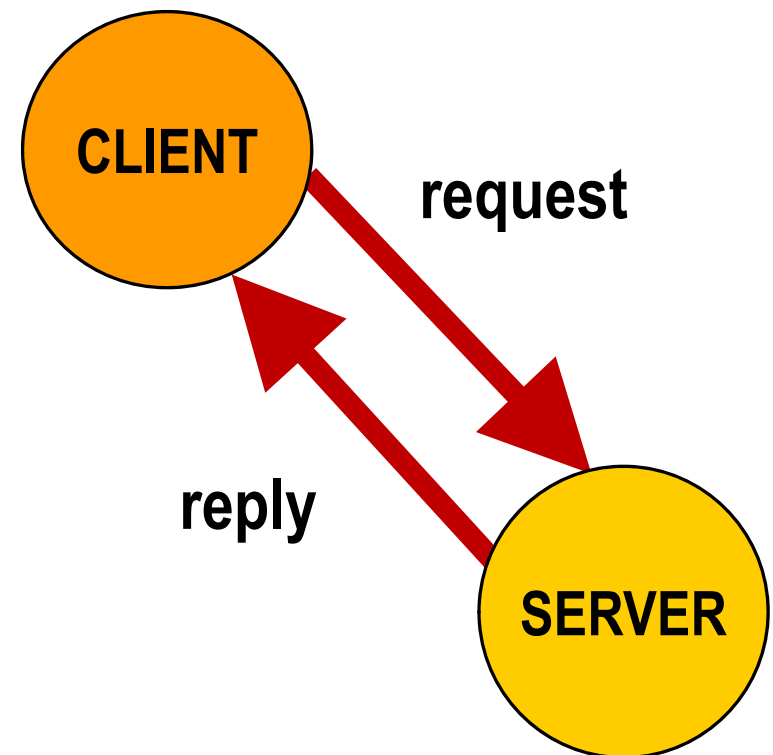


Distributed systems

- Integrated networks of computers
 - **Workstations** sharing common resources (file servers, printers, ...)
- Current trend is to leave systems very loosely coupled
 - Each computer has its own OS

Client /Server Model

- Servers wait for requests from clients and process them
 - File servers
 - Print servers
 - Authentication servers

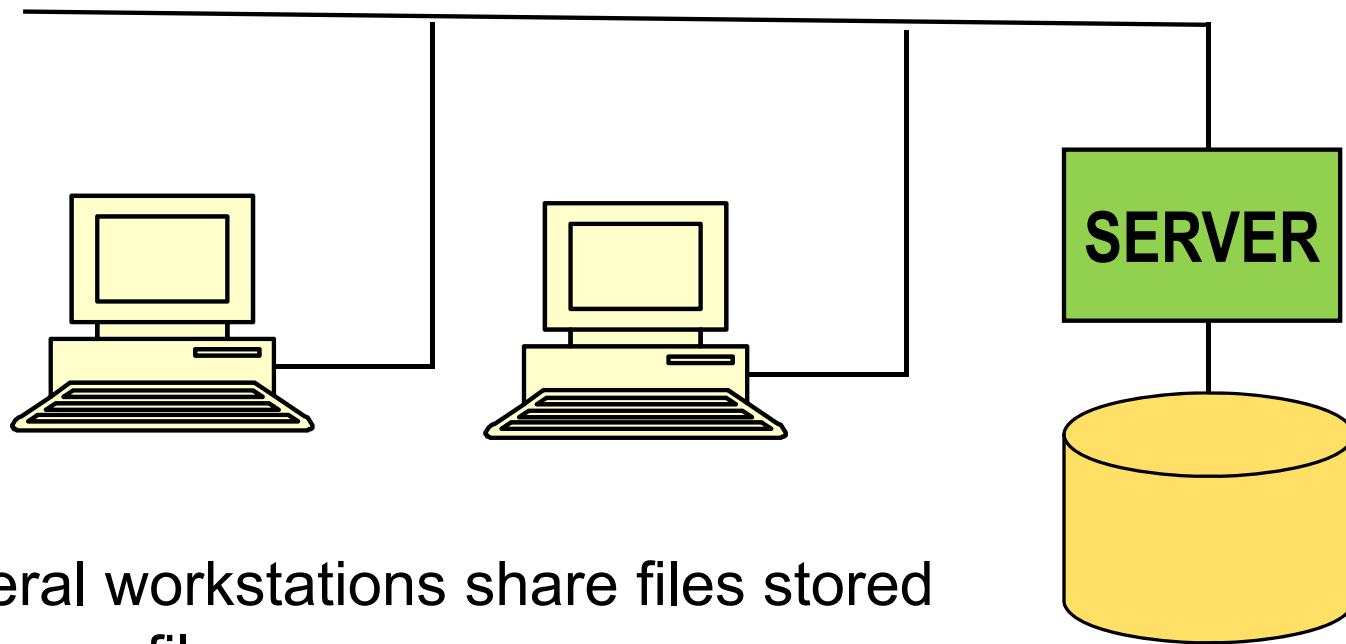




A typical sequential server

```
for (;;) {  
    // wait for request  
    get_request(...);  
    // process it  
    process_request(...);  
    // send reply  
    send_reply(...);  
} // forever
```

Network file system



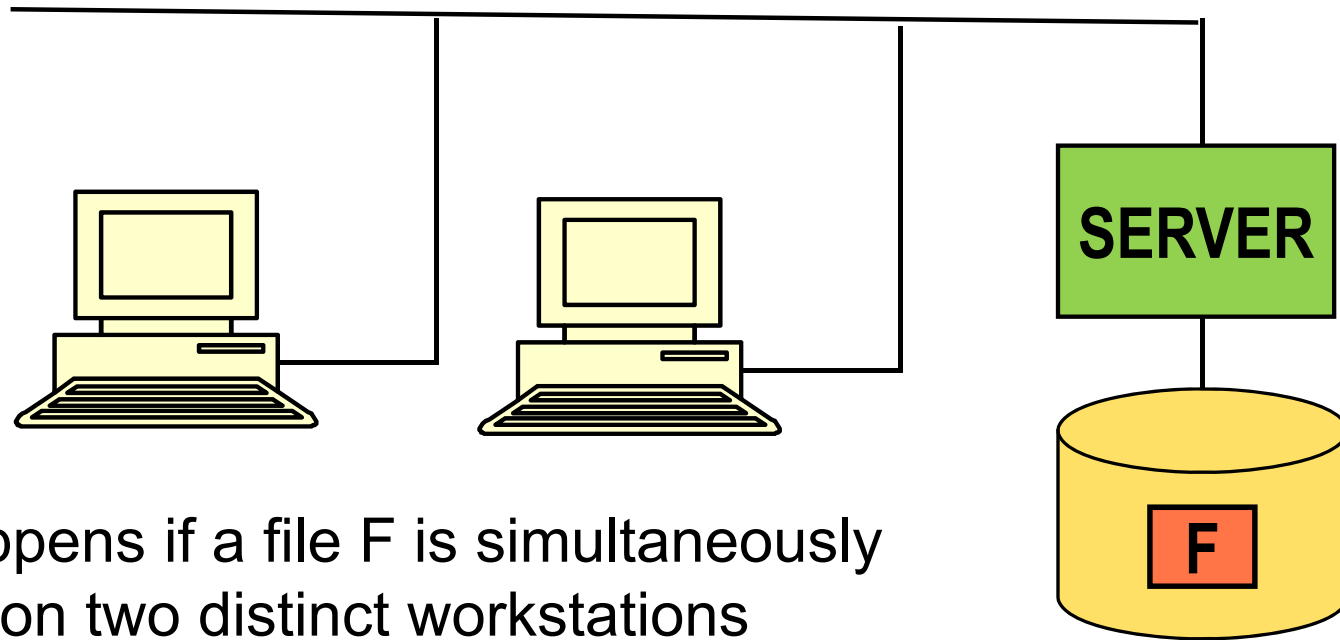
- Lets several workstations share files stored on a common file server



Performance Issues

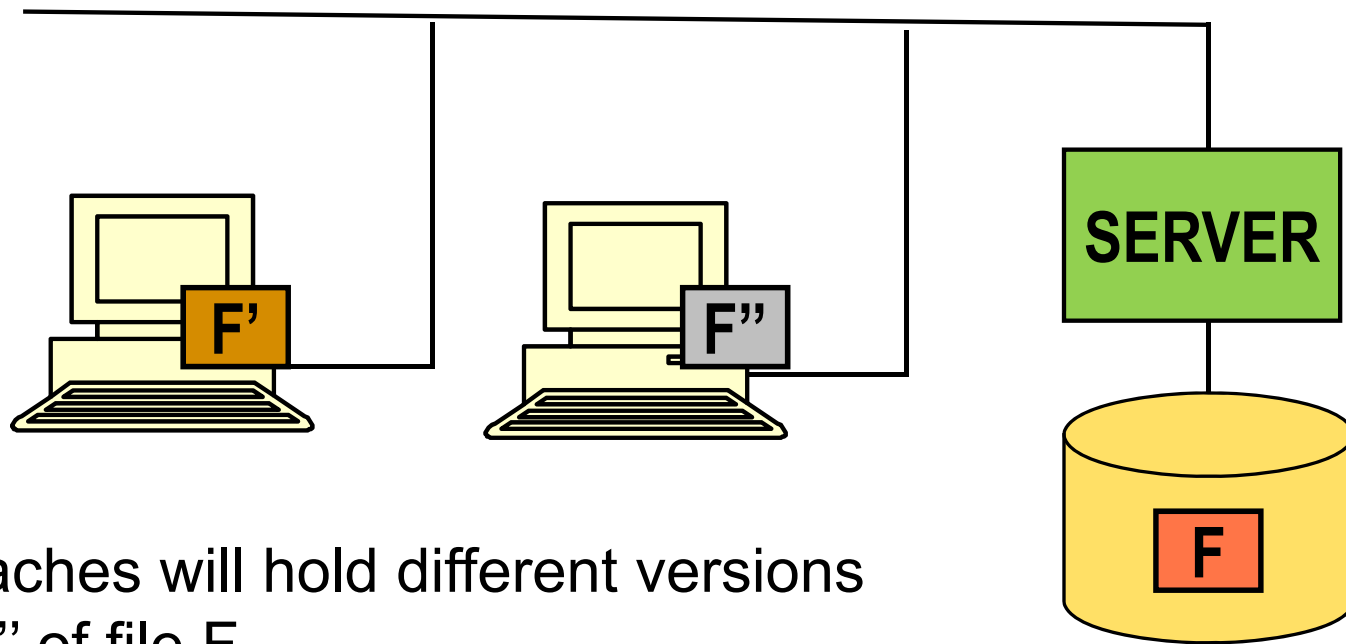
- ***Response time*** is the main issue
 - ***Network latency*** is now added to ***disk latency***
- Will attempt to mask these two latencies
 - Extensive ***client caching***
 - ***Works very well***

File consistency issues (I)



- What happens if a file F is simultaneously modified on two distinct workstations

File consistency issues (II)



- Client caches will hold different versions F' and F'' of file F



File Consistency Issues (III)

- Maintaining file consistency is a very important issue in distributed/networked file system design
- Different systems use different approaches
 - NFS from Sun Microsystems
 - AFS/Coda from CMU
 - ...



Other distributed systems issues

- **Authenticating** users
 - *A problem in open networks*
- Making distributed systems as *reliable* as stand-alone systems
 - ***Replication*** of data and services
- Keeping the clocks of the machines more or less synchronized.



Unix and Linux



UNIX (I)

- Started at Bell Labs in the early 70's as an attempt to build a sophisticated time-sharing system on a very small minicomputer.
- First OS to be almost entirely written in C
- Ported to the VAX architecture in the late 70's at U. C. Berkeley:
 - Added virtual memory and networking

The fathers of UNIX



Ken Thompson and Denis Ritchie



UNIX (II)

- Became the standard operating systems for ***workstations***
 - Selected by Sun Microsystems
- Became less popular because
 - Too many variants
 - Berkeley BSD, ATT System V, ...
 - PCs displaced workstations
 - Windows has a better user interface



UNIX Today

- Several ***free versions*** exist (FreeBSD, Linux):
 - Free access to source code
 - Ideal platform for OS research
- ***Apple OS X*** runs on the top of an updated version of BSD
- ***Android*** runs on top of a heavily customized Linux kernel
- ***Chrome*** runs on top of a vanilla Linux



A Rapid Tour

- UNIX kernel is the core of the system and handles the system calls
- UNIX has several shells: **sh**, **csch**, **ksh**, **bash**
- On-line command manual:
 - **man xyz**
displays manual page for command **xyz**
 - **man 2 xyz**
displays manual page for *system call* **xyz(...)**



Most Lasting Impact

- First OS that
 - Run efficiently on very different platforms
 - Had its source code made available to its users
- File system inspired most more recent OSes
- Remains the best platform for OS research



Kernel organizations



Kernel Organizations

- Three basic organizations:
 - ***Monolithic kernels:***
 - The default
 - ***Layered kernels:***
 - A great idea that did not work
 - ***Microkernels:***
 - Hurt by the high cost of context switches

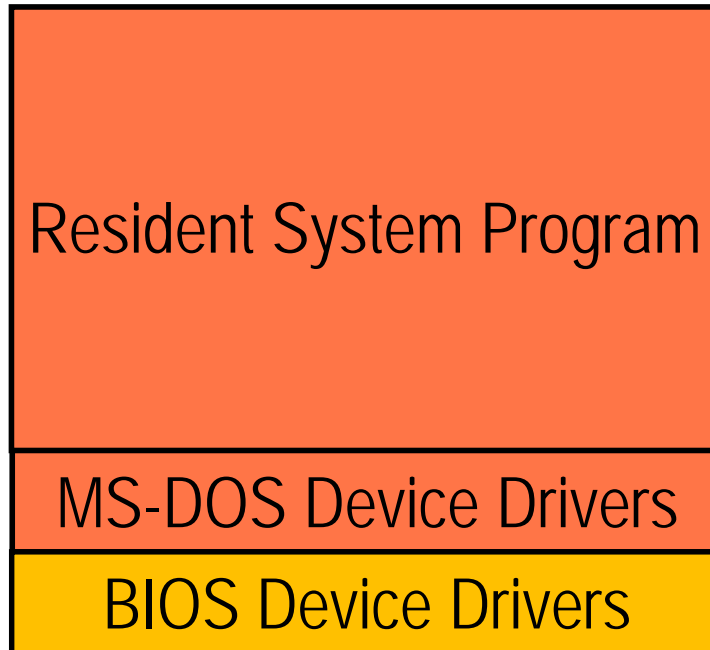


Monolithic kernels

- No particular organization
 - All kernel functions share the same address space
 - This includes ***devices drivers*** and other ***kernel extensions***
- Lack of internal organization makes the kernel ***hard to manage, extend and debug***



MSDOS (I)





The BIOS

- Basic Input-Output System
- Stored on a chip
 - First ROM, now EEPROM
- Takes control of CPU when system is turned on
 - Identifies system components
 - Initiates booting of operating system
- Also provides low-level I/O access routines



The “curse”

- Hardware lacked dual mode and hardware memory protection
 - Nothing prevented application programs from accessing directly the BIOS
 - Program accessing disk files through BIOS I/O routines assumed a given disk organization
 - Changing it became impossible



The solution

- For a long time, Microsoft could not make radical changes to its FAT-16 disk organization
- Windows XP and all modern operating systems prevent user programs from bypassing the kernel.



UNIX



Monolithic kernel

Terminal, device and memory controllers

- Monolithic kernel contains everything that is ***not device-specific*** including file system, networking code, and so forth.



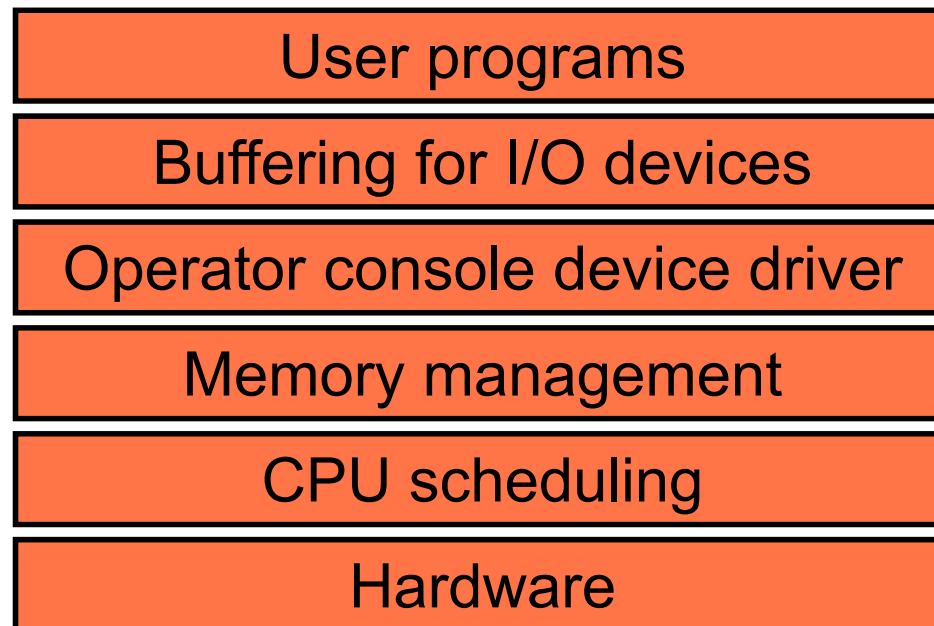
Layered kernel

- Proposed by Edsger Dijkstra
- Implemented as a hierarchy of **layers**:
- Each layer defines a new data object
 - Hiding from the higher layers some functions of the lower layers
 - Providing some new functionality



THE operating system kernel

- (named after Dutch initials of T. U. Eindhoven)





Limitations

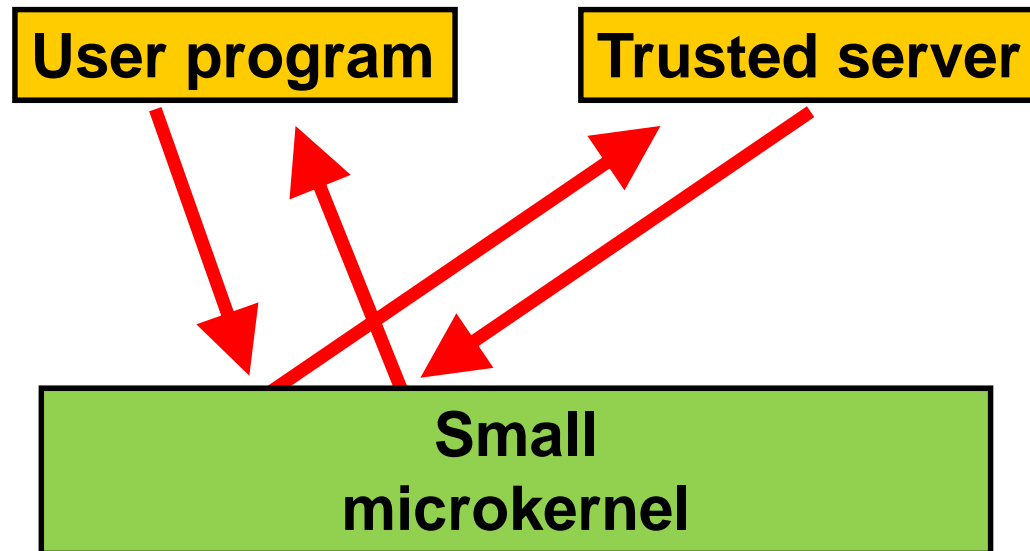
- Layered design works extremely well for networking code
 - Each layer offers its own functionality
- Much less successful for kernel design
 - No clear ordering of layers
 - Memory management uses file system features and vice versa



Microkernels

- A reaction against “bloated” monolithic kernels
 - Hard to manage, extend, debug and **secure**
- Key idea is making kernel ***smaller*** by delegating ***non-essential tasks*** to ***trusted user-level servers***
 - *Same idea as subcontracting*
- Microkernel keeps doing what cannot be delegated

How it works (I)





How it works (II)

- Microkernel
 - Receives request from user program
 - Decides to forward it to a user-level server
 - Waits for reply for server
 - Forwards it to user program
- ***Trusted servers*** run outside the kernel
 - Cannot execute privileged instructions



Advantages

- Kernel is smaller, easier to secure and manage
- Servers run outside of the kernel
 - Cannot crash the kernel
 - Much easier to extend kernel functionality
 - Adding new servers
 - Adding an NTFS server to UNIX microkernel



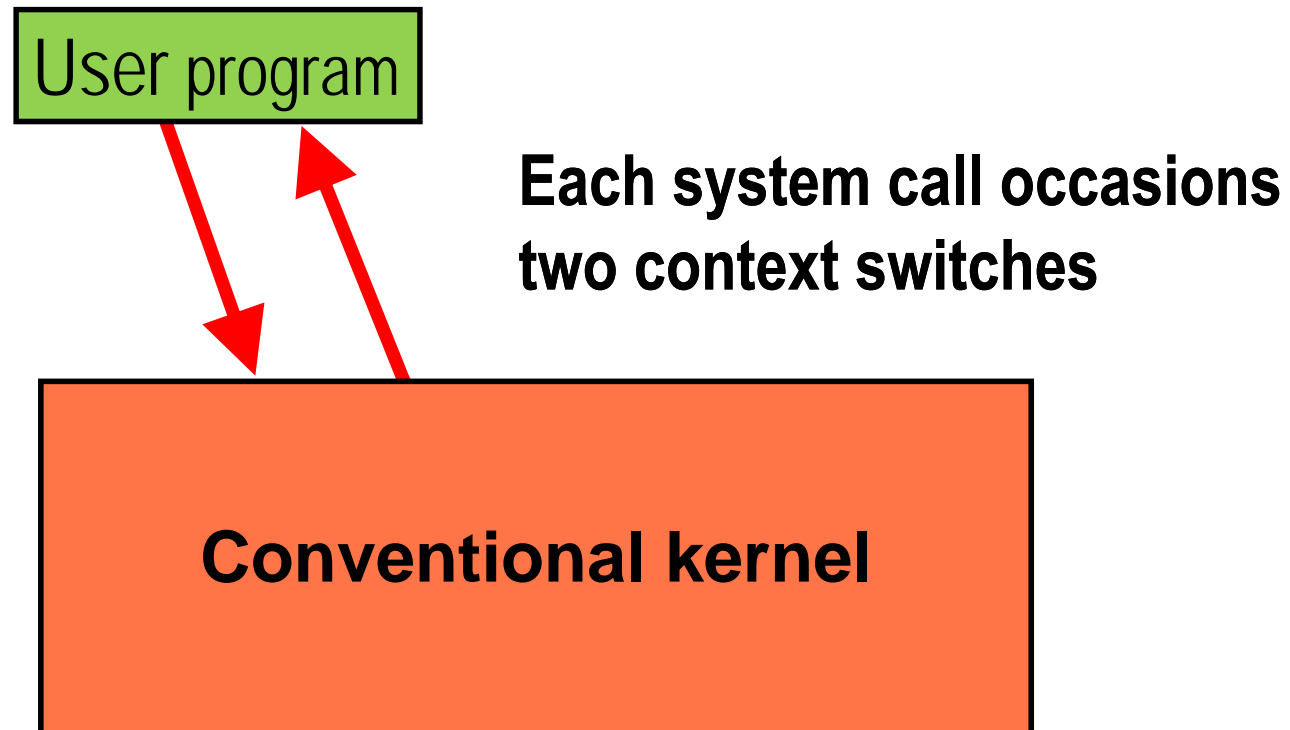
Major disadvantage

- ***Too slow***

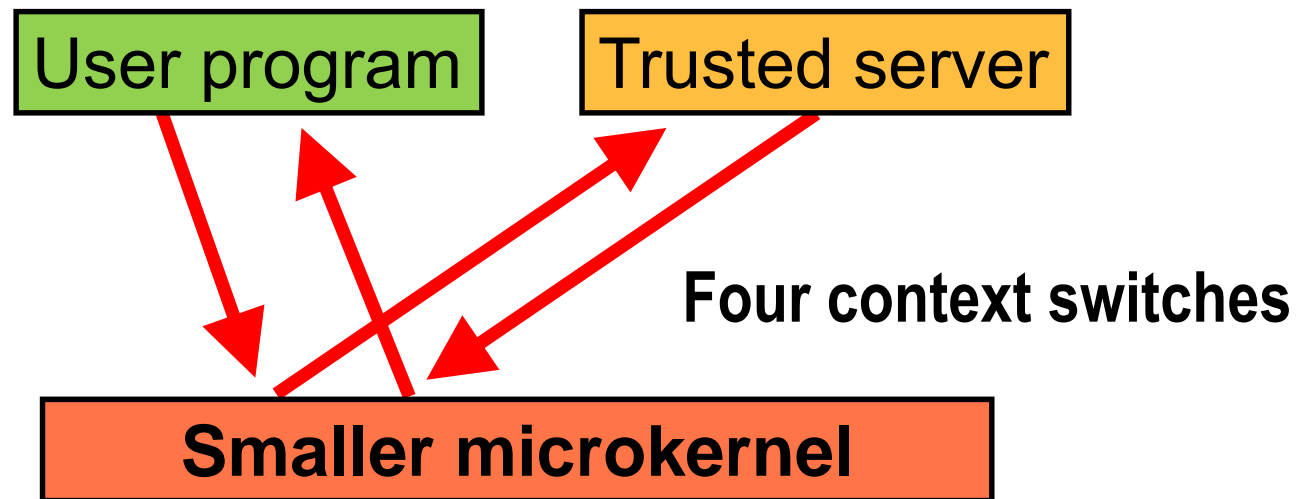
- Four context switches instead of two

- *Speed remains an essential concern*
- *We don't like to trade speed for safety (or anything else)*

A conventional kernel



A microkernel





Mach

- Designed in mid 80's to replace UNIX kernel
- New kernel with different system calls
 - UNIX system calls are routed to an ***emulation server***
- Emulation server was d to run in user space
 - Slowed down the system
 - Server ended inside the kernel



MINIX 3

- MINIX 1 was designed for teaching OS internals
 - Predates Linux
- Now aimed at high reliability (embedded) applications
 - *More willing to trade space for reliability*
- Runs on x86 and ARM processors
- Compatible with NetBSD



MINIX 3 microkernel

- "Tiny" (12,700 lines) microkernel
 - Handles ***interrupts*** and ***message*** passing
 - Only code running in kernel mode
- Other OS functions are handled by ***isolated, protected, user-mode*** processes
 - Each device driver is a separate user-mode process
 - System automatically restarts ***crashed drivers***



Modular kernels

- Linux, Windows
- Modules are object files whose contents can be linked to—and unlinked from—the kernel at any time
 - Run inside the kernel address space
 - *Used to add to the kernel **device drivers** for new devices*



Advantages

- ***Extensibility:***

- ☐ Can add new features the kernel
- ☐ In many cases, the process is completely transparent to the user

- ***Lack of performance penalty:***

- ☐ Modules run in the kernel address space



Disadvantages

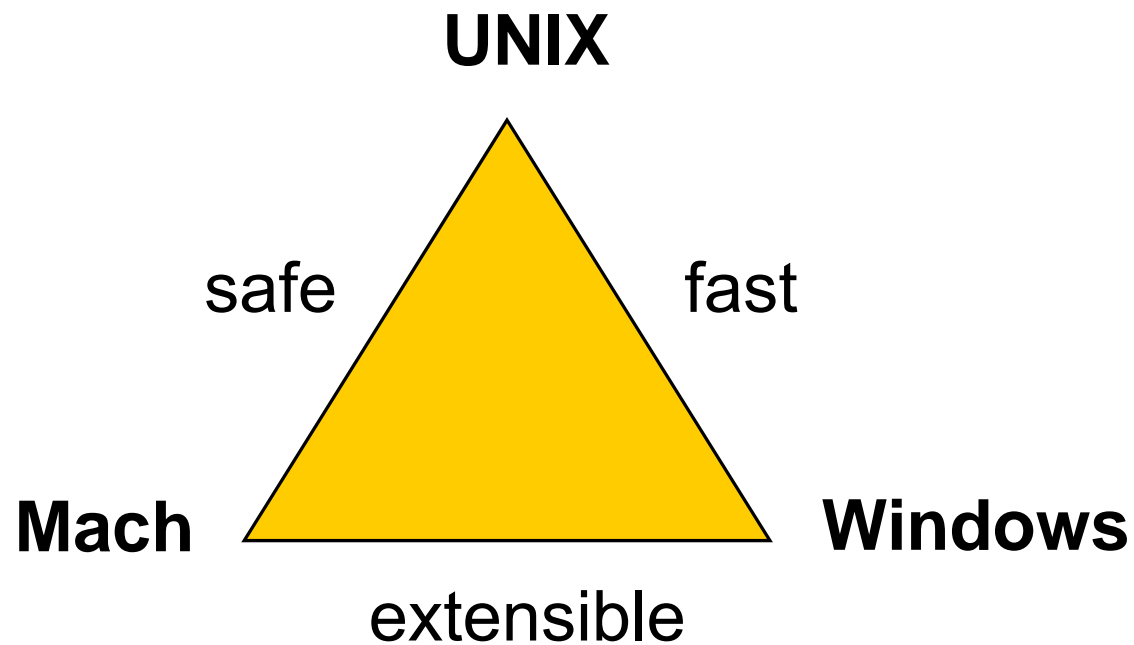
- ***Lower reliability***


- ☐ A bad module can corrupt the whole kernel and crash the system.

- ***Serious problem***

- ☐ Many device drivers are poorly written
- ☐ Device drivers account for 85% of reported failures of Windows XP

Current state of the art





Why?

- Unix has a monolithic kernel (which makes it fast) and does not allow extensions (which makes it both safe and non extensible)
- Windows has a monolithic kernel (which makes it fast) and allows extensions (which makes it both extensible and unsafe)
- Mach allows extensions in user space (which makes it extensible, safe and slow)



Virtual machines



Virtual machines

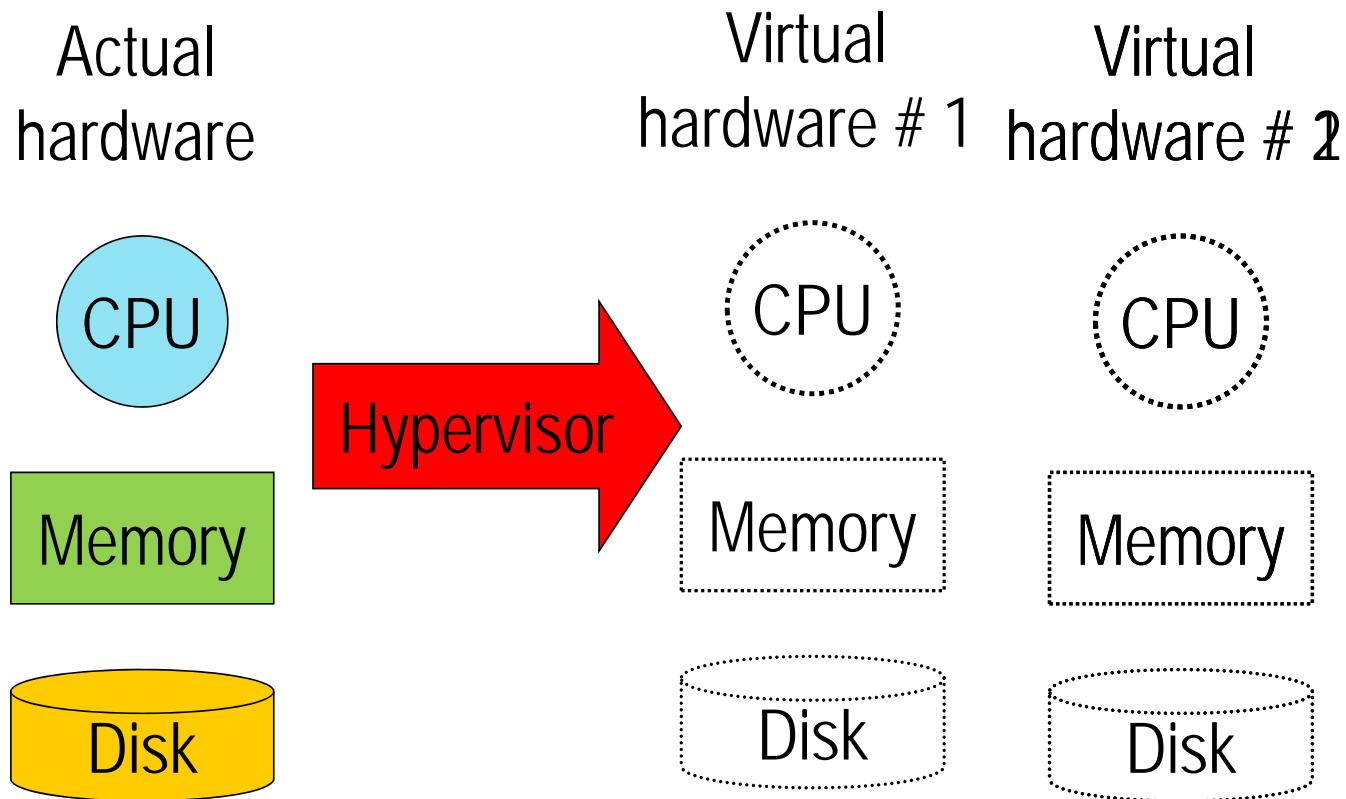
- Let different operating systems run at the same time on a single computer
 - Windows, Linux and Mac OS
 - A real-time OS and a conventional OS
 - A production OS and a new OS being tested



How it is done

- A **hypervisor / VM monitor** defines two or more virtual machines
 - Each virtual machine has
 - Its own virtual CPU
 - Its own virtual physical memory
 - Its own virtual disk(s)
- Can also install VM on top of a **host OS**
 - **VM Box**

The virtualization process

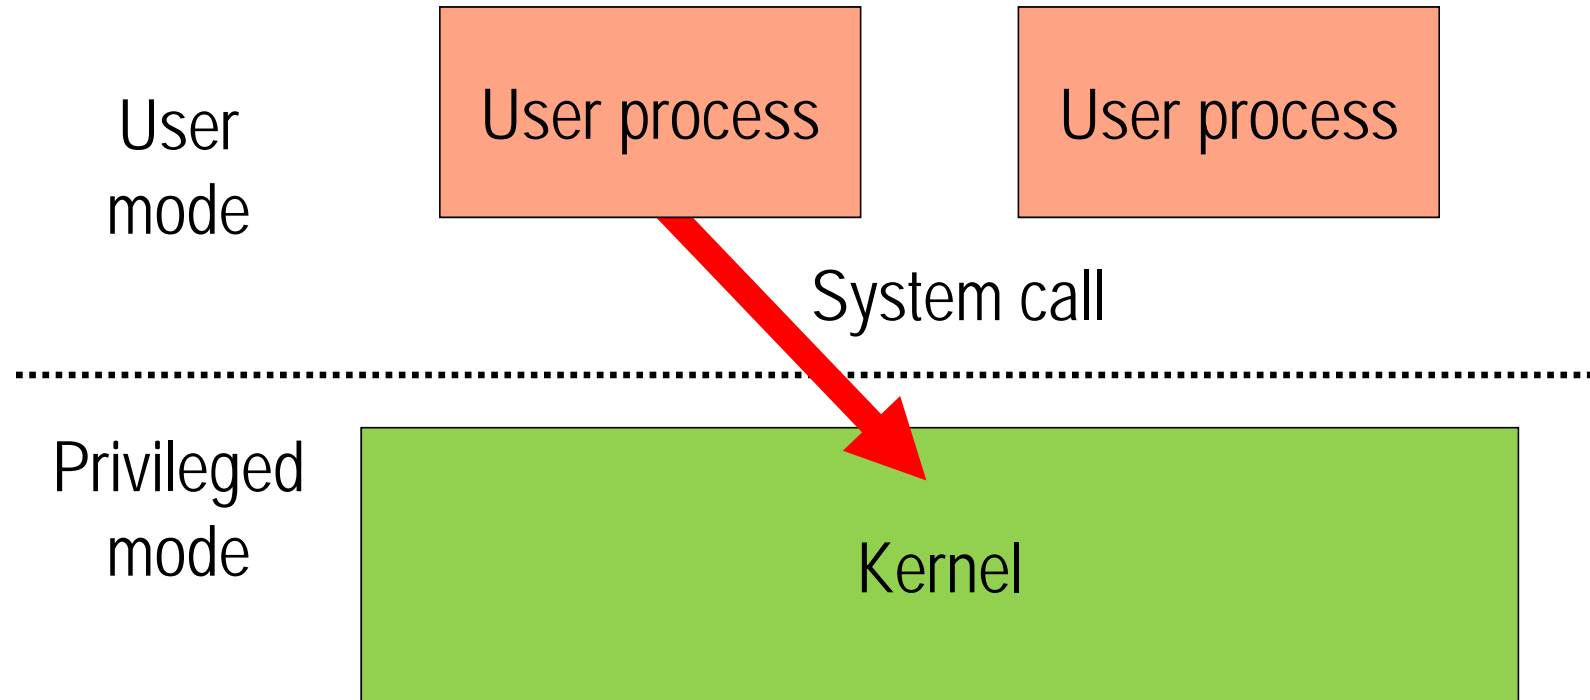




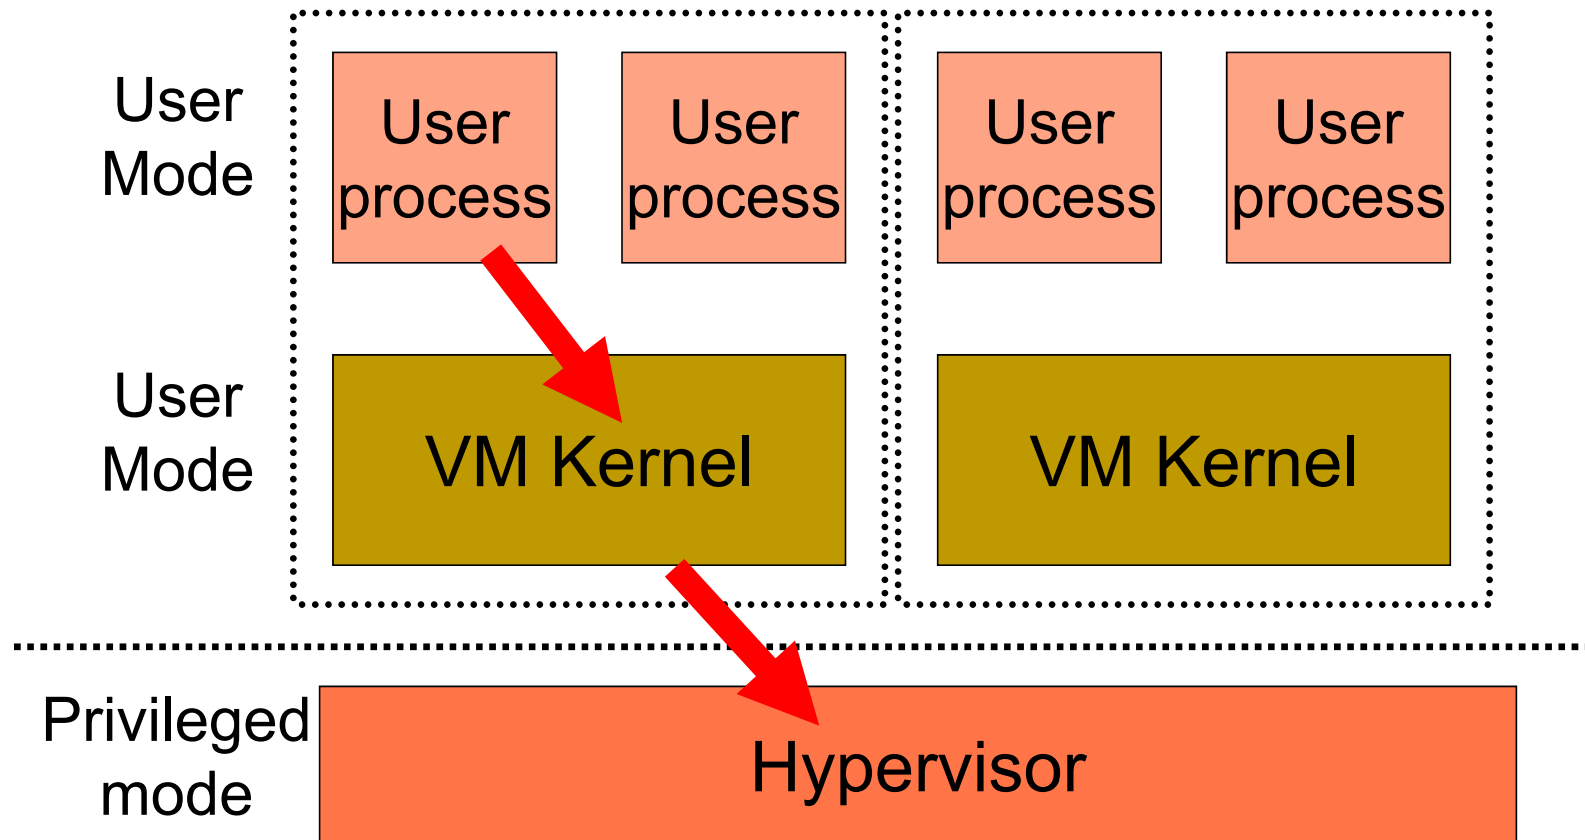
Reminder

- In a conventional OS,
 - Kernel executes in ***privileged/supervisor mode***
 - Can do virtually everything
 - User processes execute in ***user mode***
 - Cannot modify their page tables
 - Cannot execute privileged instructions

A conventional architecture



Two virtual machines

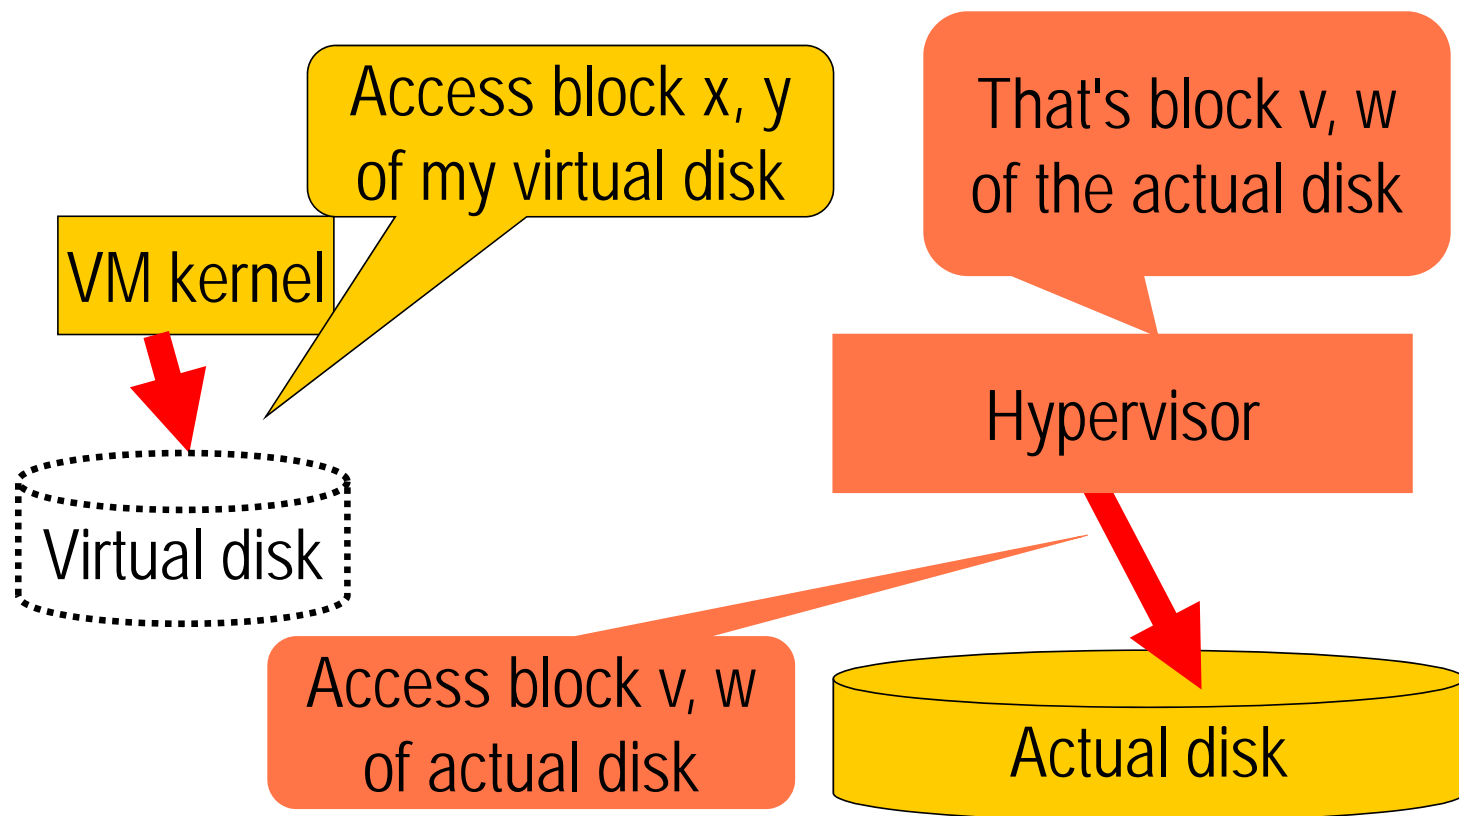




Explanations (II)

- Whenever the kernel of a VM issues a privileged instruction, an interrupt occurs
 - The hypervisor takes control and do the physical equivalent of what the VM attempted to do:
 - Must convert virtual RAM addresses into physical RAM addresses
 - Must convert virtual disk block addresses into physical block addresses

Translating a block address





Handling I/Os

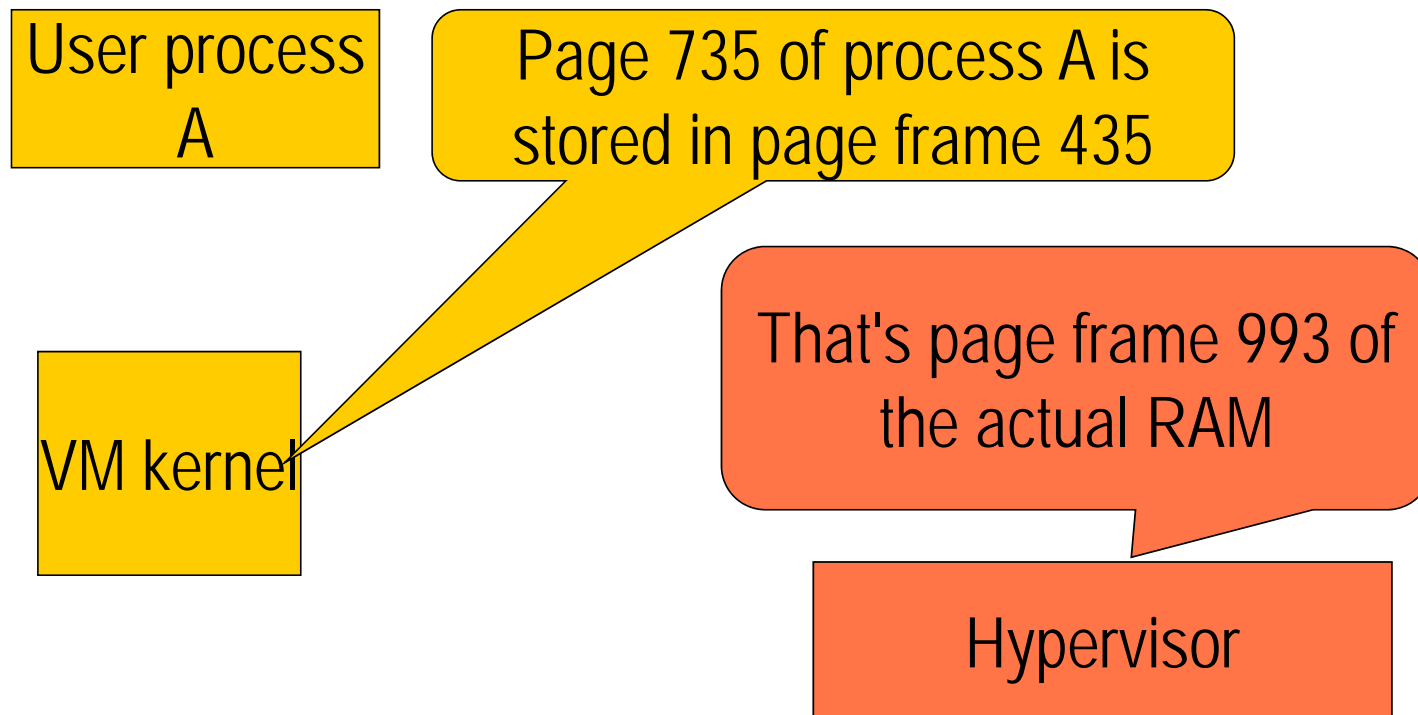
- Difficult task because
 - Wide variety of devices
 - Some devices may be shared among several VMs
 - Printers
 - Shared disk partition
 - *Want to let Linux and Windows access the same files*



Virtual Memory Issues

- Each VM kernel manages its own memory
 - Its page tables map program virtual addresses into ***what it believes to be physical addresses***

The dilemma





Nastiest Issue

- The whole VM approach assumes that a kernel executing in user mode will behave exactly like a kernel executing in privileged mode except that privileged instructions will be trapped
- ***Not true for all architectures!***
 - ***Intel x86 Pop flags (POPF) instruction***
 - ...



The Virtual Box Solution

- Code Scanning and Analysis Manager (CSAM)
 - Scans privileged code recursively before its first execution to identify problematic instructions
 - Calls the Patch Manager (PATM) to perform *in-situ* patching.



The Xen solution

- Modify the guest kernel to eliminate badly behaving instructions such as POPF
 - ***Paravirtualization***
 - Faster but less flexible
 - Requires open-source kernel

User programs are not affected

- Only the kernel



Containers

- Each VM runs its own copy of the kernel
 - Takes memory space
- Containers provide isolated user-space instances that share the same kernel
 - Less overhead
 - Less flexibility
- Docker, LYXC