



COSC 3360/6310

Wednesday, February 3

Announcements



- A ***sample input*** for the first assignment is now on MS Teams
 - **input10.txt**
 - Checked for correctness
- ***First quiz*** on ***February 8*** at ***4pm*** on ***Blackboard***
 - A ***practice quiz*** will be ***online today at 7pm***
 - Will be remain available until Monday afternoon
 - Please mail me (jfparis@uh.edu) any issues you might have

Materials on the quiz



- Anything discussed in class, except assignment hints:
 - **A. Wednesday January 20.pptx**
 - **B. Monday January 25.pptx**
 - **C. Wednesday January 27.pptx**
 - **D. Monday February 1.pptx**
 - **E. Wednesday February 3.pptx**



Chapter I

Introduction

Jehan-François Pâris
jfparis@uh.edu



Unix and Linux

UNIX (I)



- Started at Bell Labs in the early 70's as an attempt to build a sophisticated time-sharing system on a very small minicomputer.
- First OS to be almost entirely written in C
- Ported to the VAX architecture in the late 70's at U. C. Berkeley:
 - Added virtual memory and networking

The fathers of UNIX



Ken Thompson and Denis Ritchie

UNIX (II)



- Became the standard operating systems for ***workstations***
 - Selected by Sun Microsystems
- Became less popular because
 - Too many variants
 - Berkeley BSD, ATT System V, ...
 - PCs displaced workstations
 - Windows has a better user interface

UNIX Today



- Several **free versions** exist (FreeBSD, Linux):
 - Free access to source code
 - Ideal platform for OS research
- **Apple OS X** runs on the top of an updated version of BSD
- **Android** runs on top of a heavily customized Linux kernel
- **Chrome** runs on top of a vanilla Linux OS

A Rapid Tour



- UNIX kernel is the core of the system and handles the system calls
- UNIX has several shells: **sh**, **cs****h**, **ksh**, **bash**
- On-line command manual:
 - **man xyz**
displays manual page for command **xyz**
 - **man 2 xyz**
displays manual page for *system call* **xyz(...)**

Most Lasting Impact



- First OS that
 - Run efficiently on very different platforms
 - Had its source code made available to its users
- File system inspired most more recent OSes
- Remains the best platform for OS research



Kernel organizations

Kernel Organizations



- Three basic organizations:
 - ***Monolithic kernels:***
 - The default
 - ***Layered kernels:***
 - A great idea that did not work
 - ***Microkernels:***
 - Hurt by the high cost of context switches

Monolithic kernels



- No particular organization
 - All kernel functions share the same address space
 - This includes ***devices drivers*** and other ***kernel extensions***
- Lack of internal organization makes the kernel ***hard to manage, extend and debug***

MS-DOS (I)



Resident System Program

MS-DOS Device Drivers

BIOS Device Drivers



The BIOS



- Basic Input-Output System
- Stored on a chip
 - First ROM, now EEPROM
- Takes control of CPU when system is turned on
 - Identifies system components
 - Initiates booting of operating system
- Also provides low-level I/O access routines

The “curse”



- Hardware lacked dual mode and hardware memory protection
 - Nothing prevented application programs from accessing directly the BIOS
 - Program accessing disk files through BIOS I/O routines assumed a given disk organization
 - Changing it became impossible



The solution



- For a long time, Microsoft could not make radical changes to its FAT-16 disk organization
- Windows XP and all modern operating systems prevent user programs from bypassing the kernel.

UNIX



Monolithic kernel

Terminal, device and memory controllers

- Monolithic kernel contains everything that is ***not device-specific*** including file system, networking code, and so forth.

Layered kernel

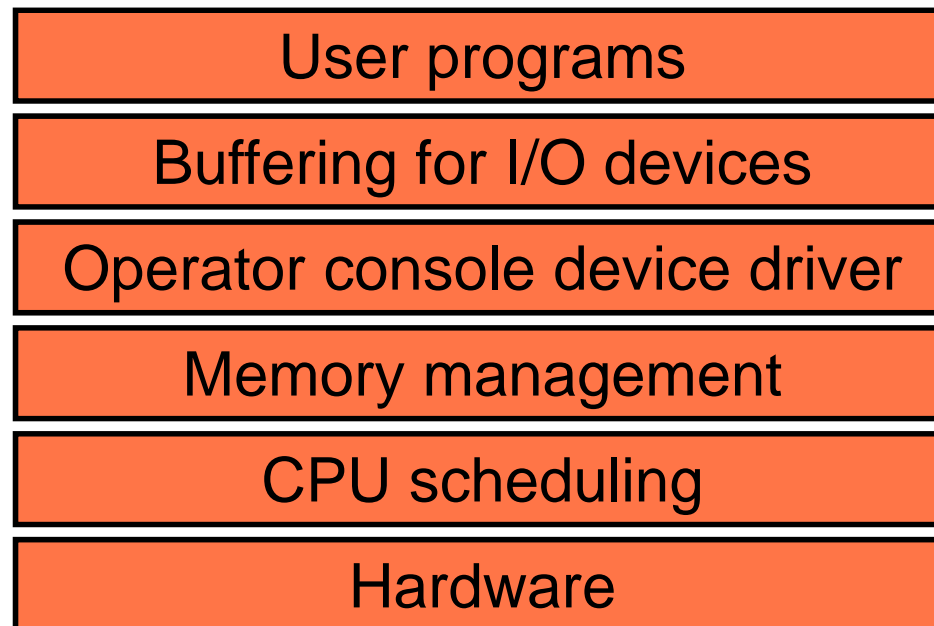


- Proposed by Edsger Dijkstra
- Implemented as a hierarchy of **layers**:
- Each layer defines a new data object
 - Hiding from the higher layers some functions of the lower layers
 - Providing some new functionality

THE operating system kernel



- (named after Dutch initials of T. U. Eindhoven)



Limitations



- Layered design works extremely well for ***networking code***
 - Each layer offers its own functionality
- Much less successful for kernel design
 - No clear ordering of layers
 - Memory management uses file system features and vice versa




Microkernels



- A reaction against “bloated” monolithic kernels
 - Hard to manage, extend, debug and **secure**
- Key idea is making kernel ***smaller*** by delegating ***non-essential tasks*** to ***trusted user-level servers***
 - *Same idea as subcontracting*
- Microkernel keeps doing what cannot be delegated:
 - Security, short-term scheduling, ...

How it works (I)



How it works (II)



- Microkernel
 - Receives request from user program
 - Decides to forward it to a user-level server
 - Waits for reply for server
 - Forwards it to user program
- ***Trusted servers*** run outside the kernel
 - Cannot execute privileged instructions



Advantages



- Kernel is smaller, easier to secure and manage
- Servers run outside of the kernel
 - Cannot crash the kernel
 - Much easier to extend kernel functionality
 - Adding new servers
 - Adding an NTFS server to UNIX microkernel

Major disadvantage

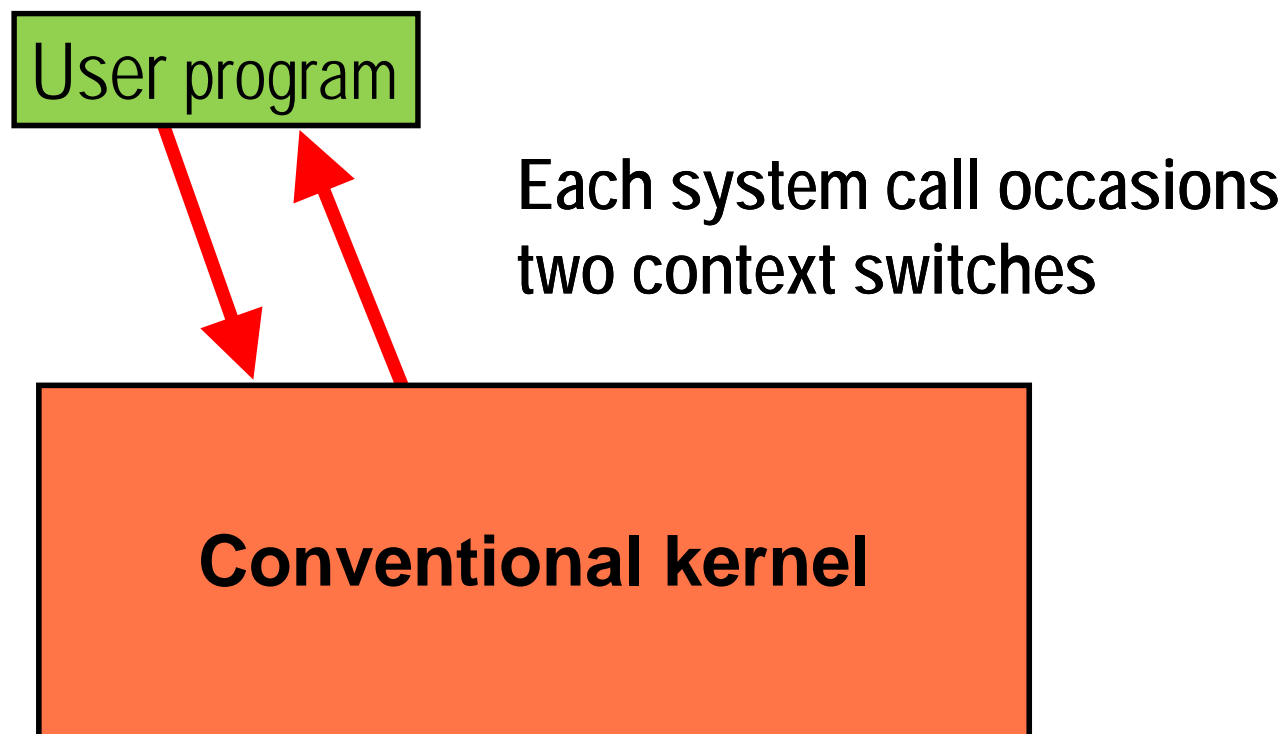


- ***Too slow***


- Four context switches instead of two

- *Speed remains an essential concern*
- *We don't like to trade speed for safety (or anything else)*

A conventional kernel



A microkernel



Mach



- Designed in mid 80's to replace UNIX kernel
- New kernel with different system calls
 - UNIX system calls are routed to an ***emulation server***
- Emulation server was d to run in user space
 - Slowed down the system
 - Server ended inside the kernel

MINIX 3



- MINIX 1 was designed for teaching OS internals
 - Predates Linux
- Now aimed at high reliability (embedded) applications
 - *More willing to trade space for reliability*
- Runs on x86 and ARM processors
- Compatible with NetBSD

MINIX 3 microkernel



- "Tiny" (12,700 lines) microkernel
 - Handles ***interrupts*** and ***message*** passing
 - Only code running in kernel mode
- Other OS functions are handled by ***isolated, protected, user-mode*** processes
 - Each device driver is a separate user-mode process
 - System automatically restarts ***crashed drivers***

Modular kernels



- Linux, Windows
- Modules are object files whose contents can be linked to—and unlinked from—the kernel at any time
 - Run inside the kernel address space
 - *Used to add to the kernel **device drivers** for new devices*

Advantages



- ***Extensibility:***

- ☐ Can add new features the kernel
- ☐ In many cases, the process is completely transparent to the user

- ***Lack of performance penalty:***

- ☐ Modules run in the kernel address space

Disadvantages




- ***Lower reliability***

- ☐ A bad module can corrupt the whole kernel and crash the system.

- ***Serious problem***

- ☐ Many device drivers are poorly written
- ☐ Device drivers account for 85% of reported failures of Windows XP

Current state of the art



Why?



- Unix has a monolithic kernel (which makes it fast) and does not allow extensions (which makes it both safe and non-extensible)
- Windows has a monolithic kernel (which makes it fast) and allows extensions (which makes it both extensible and unsafe)
- Mach allows extensions in user space (which makes it extensible, safe and slow)



Virtual machines

Virtual machines




- Let different operating systems run at the same time on a single computer
 - Windows, Linux and Mac OS
 - A real-time OS and a conventional OS
 - A production OS and a new OS being tested

How it is done



- A **hypervisor / VM monitor** defines two or more virtual machines
 - Each virtual machine has
 - Its own virtual CPU
 - Its own virtual physical memory
 - Its own virtual disk(s)
- Can also install VM on top of a **host OS**
 - **VM Box**

The virtualization process




Reminder

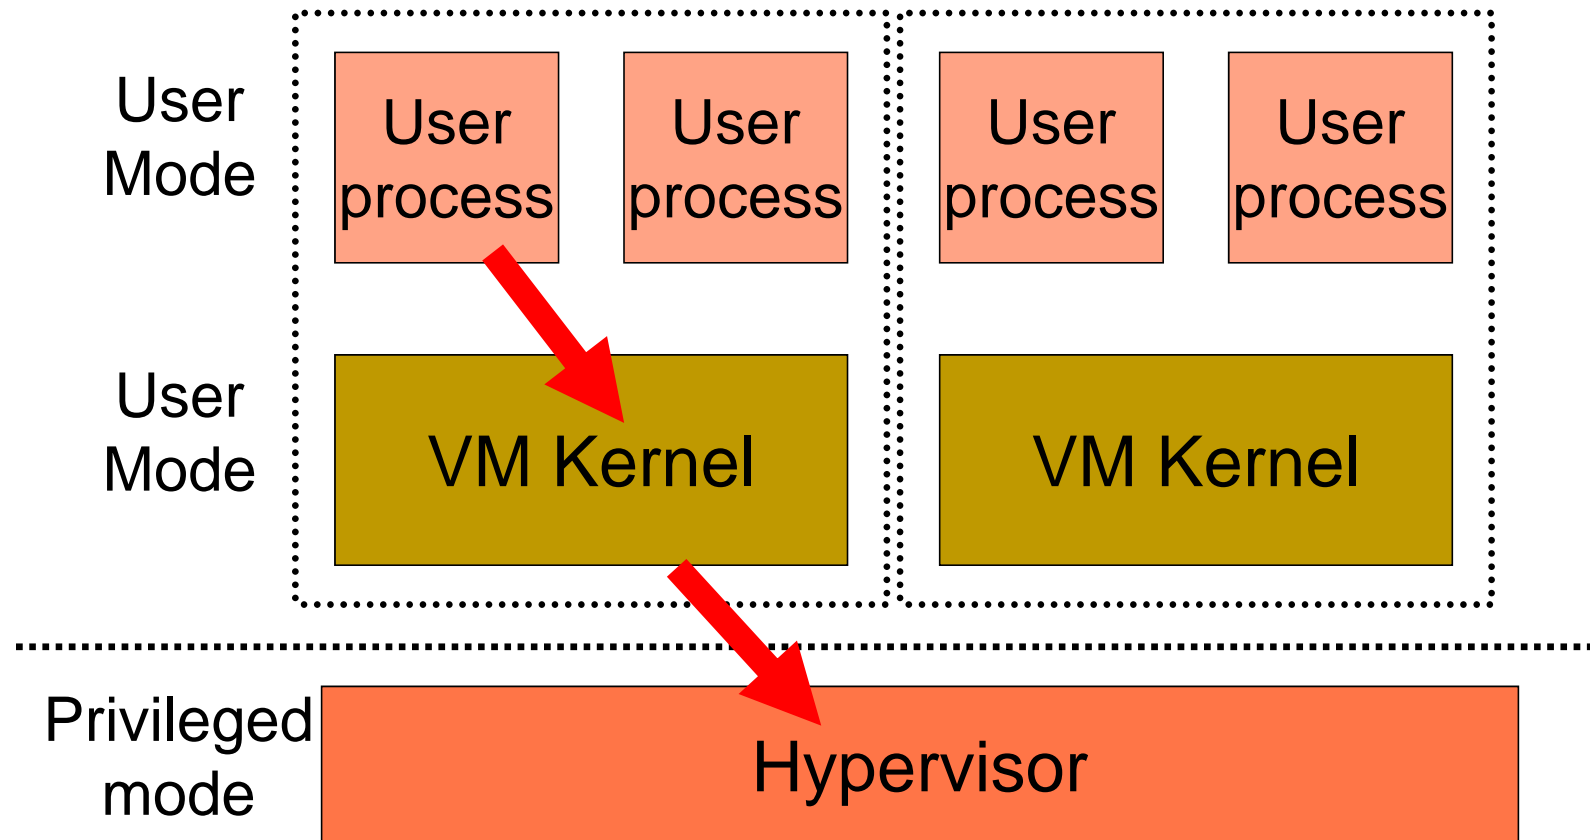


- In a conventional OS,
 - Kernel executes in ***privileged/supervisor mode***
 - Can do virtually everything
 - User processes execute in ***user mode***
 - Cannot modify their page tables
 - Cannot execute privileged instructions

A conventional architecture



Two virtual machines

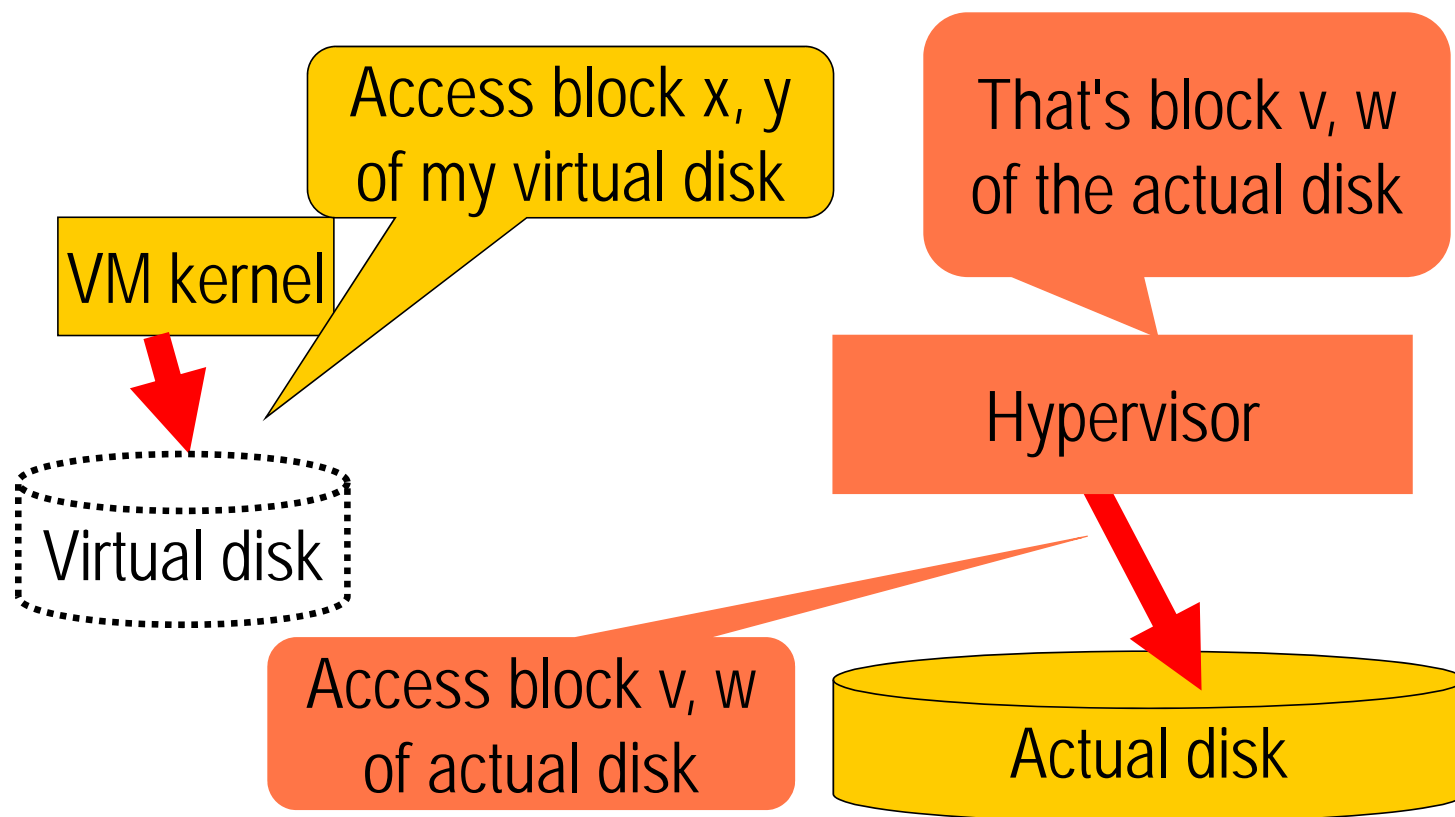


Explanations (II)



- Whenever the kernel of a VM issues a privileged instruction, an interrupt occurs
 - The hypervisor takes control and do the physical equivalent of what the VM attempted to do:
 - Must convert virtual RAM addresses into physical RAM addresses
 - Must convert virtual disk block addresses into physical block addresses

Translating a block address



Handling I/Os



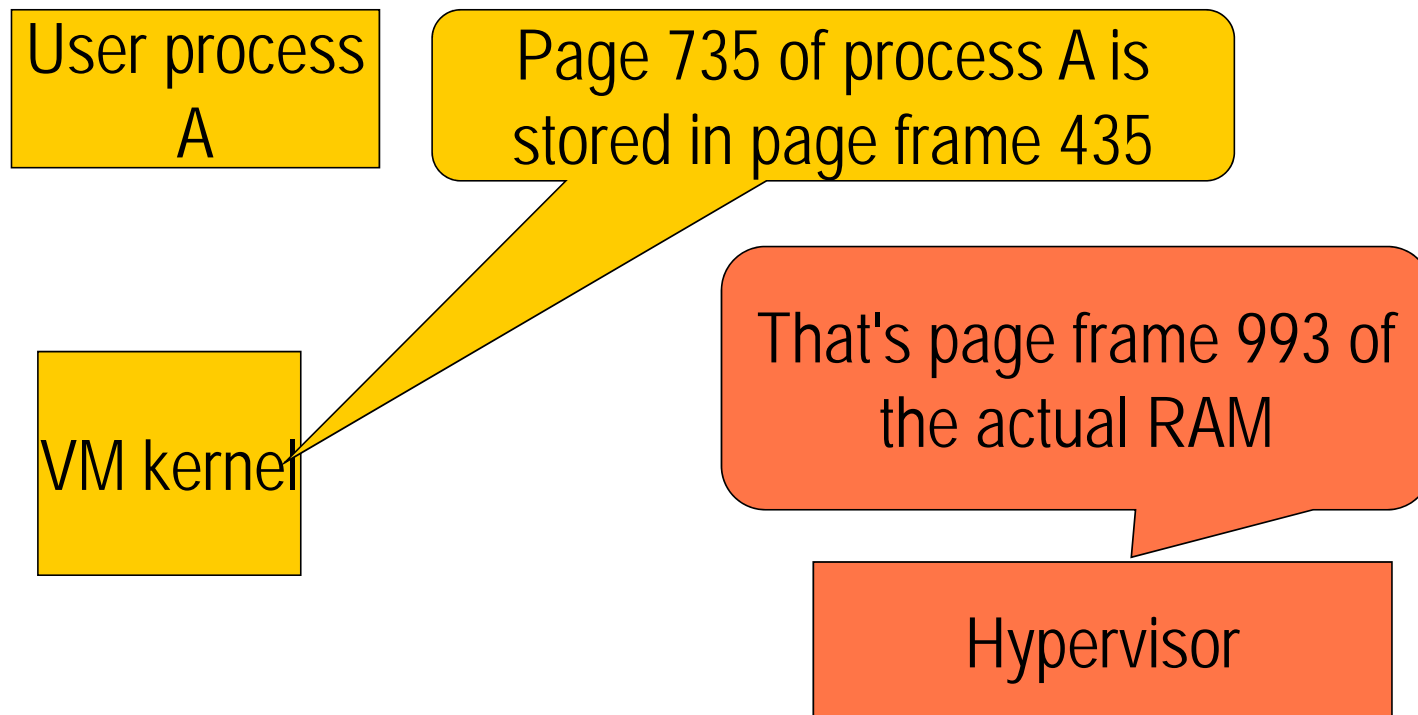
- Difficult task because
 - Wide variety of devices
 - Some devices may be shared among several VMs
 - Printers
 - Shared disk partition
 - *Want to let Linux and Windows access the same files*

Virtual Memory Issues



- Each VM kernel manages its own memory
 - Its page tables map program virtual addresses into ***what it believes to be physical addresses***

The dilemma



Nastiest Issue



- The whole VM approach assumes that a kernel executing in user mode will behave exactly like a kernel executing in privileged mode except that privileged instructions will be trapped
- ***Not true for all architectures!***
 - ***Intel x86 Pop flags (POPF) instruction***
 - ...

The Virtual Box Solution



- Code Scanning and Analysis Manager (CSAM)
 - Scans privileged code recursively before its first execution to identify problematic instructions
 - Calls the Patch Manager (PATM) to perform *in-situ* patching.

The Xen solution



- Modify the guest kernel to eliminate badly behaving instructions such as POPF
 - ***Paravirtualization***
 - Faster but less flexible
 - Requires open-source kernel

User programs are not affected
❖ **Only the kernel**

Containers



- Each VM runs its own copy of the kernel
 - Takes memory space
- Containers provide isolated user-space instances that share the same kernel
 - Less overhead
 - Less flexibility
- Docker, LYXC