




COSC 3360/6310

Wednesday, February 10


Announcements



- Please try to start your next quizzes on time. I am most attentive to issues during the first 15 to 20 minutes of the quiz.
- I plan to have all 120 quizzes graded by next Monday.
- I have extra explanations for the first assignment.




How to load jobs into main memory




- Must be done
 - At beginning of the simulation
 - Each time a job terminates
- Two conditions
 - Less than MPL jobs in main memory
 - and**
 - There still are jobs ready to be fetched

Fetchjobs routine



```
fetchjobs(seqNo) {
    while ((njobsinram < mpl)&&(nextjobtofetch < jobcount)) {
        njobsinram ++;
        seqno = nextjobtofetch;
        nextjobtofetch++;
        pop first job step from job seqno
        split into (request, duration)
        if (request == "CORE")
            process core request for job jobID[seqno]
        else
            printf("PANIC: FIRST STEP IS NOT A CORE REQUEST");
    } // fetchjobs
```

Sequence numbers and job IDs




- JobIDs are not consecutive
 - Cannot use them to locate the next step of a specific job

- My solution
 - Use job sequence number inside your program
 - Keep track of job IDs in a jobID array indexed by job sequence numbers



Overview (I)



■ ***Input module***


- Read in all input data
 - Store them in a jobList
- Fetch up to MPL jobs (**fetchjobs**)

■ ***Main loop***

- Pops next event from event list
 - CORE completion
 - DISK completion
 - SPOOLER completion



Overview (II)



- ***Starting a job***
 - Handled by `fetchjobs` routine
- ***Handling job termination***
 - Try to fetch a new job (`fetchjobs`)
- ***Detecting the end of the simulation***
 - Event list will be *empty*




Chapter II

Processes

Jehan-François Pâris
jfparis@uh.edu




Chapter Overview




- Processes
- States of a process
- Operations on processes
 - **fork()**, **exec()**, **kill()**, **signal()**
- Threads and lightweight processes

PROCESSES




- A process is a ***program*** executing a given ***sequential computation***.
 - An ***active entity*** unlike a program
 - *Think of the difference between a recipe in a cookbook and the activity of a cook preparing a dish according to the recipe!*

Processes and programs (I)



- Can have one program and many processes
 - When several users execute the same program (text editor, compiler, and so forth) at the same time, each execution of the program constitutes a ***separate process***
 - A program that ***forks*** another sequential computation gives birth to a new process.

Examples




- Several executions of same program

gcc


gcc

gcc

- A program forking a child




Processes and programs (II)




- Can have one process and two—or more—programs
 - A process that performs an `exec()` call replaces the program it was executing


Examples



- One process executing two programs




The UNIX shell



- Program that
 - Reads input from the keyboard
 - Creates the process that will execute the command.
 - Wait for the completion of the process it has created unless it was specified otherwise
- *User-level program that you and I could write*


Yes, we can



```
#!/usr/bin/python3
""" A very very basic shell in Python 3
    Check https://www.python-course.eu/forking.php
"""
import os
def changeDirectory(argc, argv) :
    if argc == 2 :
        try :
            os.chdir(argv[1])
        except Exception :
            print("BasicShell: " + argv[0] +
                  ": no such file or directory")
    elif argc == 1 :
        os.chdir(os.environ['HOME'])
    else :
        print("BasicShell: cd: too many arguments")
def vanillaCase(argc, argv) :
    kidpid = os.fork()
    if kidpid == 0 :
        try :
            os.execvp(argv[0], argv)
        except Exception :
            print(argv[0]+": program not found")
    else :
        os.wait()
```


```
while (1) :
    argline = input("BasicShell: ")
    argline.strip()
    argv = argline.split() # break at spaces
    argc = len(argv)
    if argc == 0 :
        continue
    if argv[0] == 'exit' :
        # Exiting BasicShell
        break
    elif argv[0] == 'cd':
        # Changing current directory
        changeDirectory(argc, argv)
    else :
        vanillaCase(argc, argv)
```


A very basic UNIX shell




```
■ for (;;) {  
    parse_input_line(arg_vector);  
    if built_in_command(arg_vector[0]) {  
        do_it(arg_vector);  
        continue;  
    } // built-in command  
    pathname = find_path(arg_vector[0]);  
    create_process(pathname, arg_vector);  
    if (interactive())  
        wait_for_this_child();  
} // for loop
```

Notes



- All functions in italics are templates yet to be written
- Real shells do more:
 - I/O redirection
 - Pipes (as in **ls -alg | more**)
 - Command aliasing,
 - Wildcard characters (as "*****")
 - ...


Importance of processes



- Processes are the ***basic entities*** managed by the operating system
 - OS provides to each process the illusion it has the whole machine for itself
 - Each process has a dedicated ***address space***




The process address space



- Set of main memory locations allocated to the process
 - Other processes cannot access them
 - Process cannot access address spaces of other processes
- A process address space is the ***playpen*** or the ***sandbox*** of its owner




A last word



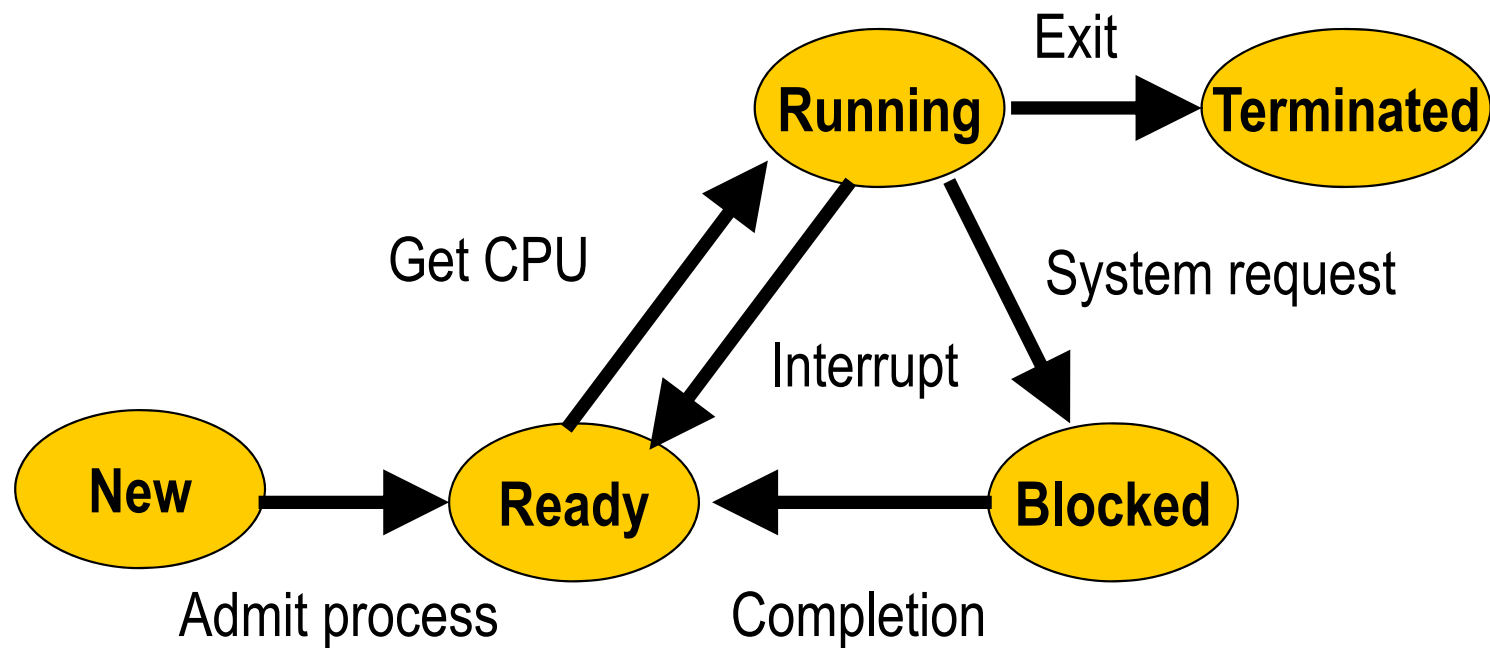

- There are many *quasi-synonyms* for process:
 - Job (very old programmers still use it)
 - Task
 - Program (strongly deprecated)

PROCESS STATES



- Processes go repeatedly through several stages during their execution
 - Waiting to get into main memory
 - Waiting for the CPU
 - Running
 - Waiting for the completion of a system call


The big diagram



This is fundamental material




Process arrival




- New process
 - Starts in NEW state
 - Gets allocated a Process Control Block (PCB) and main memory
 - Is put in the READY state waiting for CPU time

The ready state




- AKA the ***ready queue***
- Contains all processes waiting for the CPU
- Organized as a ***priority queue***
- Processes leave the priority queue when they get some CPU time
 - Move then to the RUNNING state

The running state (I)




- A process in the running state has exclusive use of the CPU until
 - It *terminates* and goes to the **TERMINATED** state
 - It does a *system call* and goes to the **BLOCKED** state
 - It is *interrupted* and returns to the **READY** state

The running state (II)




- Processes are forced to relinquish the CPU and return to the **READY** state when
 - A **higher-priority process** arrives in the ready queue and **preempts** the running process
 - *Get out, I'm more urgent than you!*
 - A **timer interrupt** indicates that the process has exceeded its time slice of CPU time

The blocked state (I)




- Contains all processes waiting for the completion of a system request:
 - I/O operation
 - Any other system call
- Process is said to be
 - ***blocked*** (Arpaci-Dusseau & Arpaci-Dusseau)
 - ***waiting***
 - ***sleeping*** (UNIX)

The blocked state (II)




- A system call that does not require callers to wait until its completion is said to be ***non-blocking***
 - Calling processes are immediately returned to the ***READY*** state
- The blocked state is organized as a ***set of queues***
 - One queue per device, OS resource

The process control block (I)




- Contains all the information associated with a specific process:
 - **Process identification** (pid), *argument vector*, ...
 - UNIX pids are unique integers
 - **Process state** (new, ready, running, ...),
 - **CPU scheduling information**
 - Process priority, processors on which the process can run, ...,

The process control block (II)




- ***Program counter*** and other CPU registers
 - Including the ***Program Status Word*** (PSW),
- ***Memory management information***
 - Very system specific,
- ***Accounting information***
 - CPU time used, system time used, ...
- ***I/O status information***
 - List of opened files, allocated devices, ...

The process table



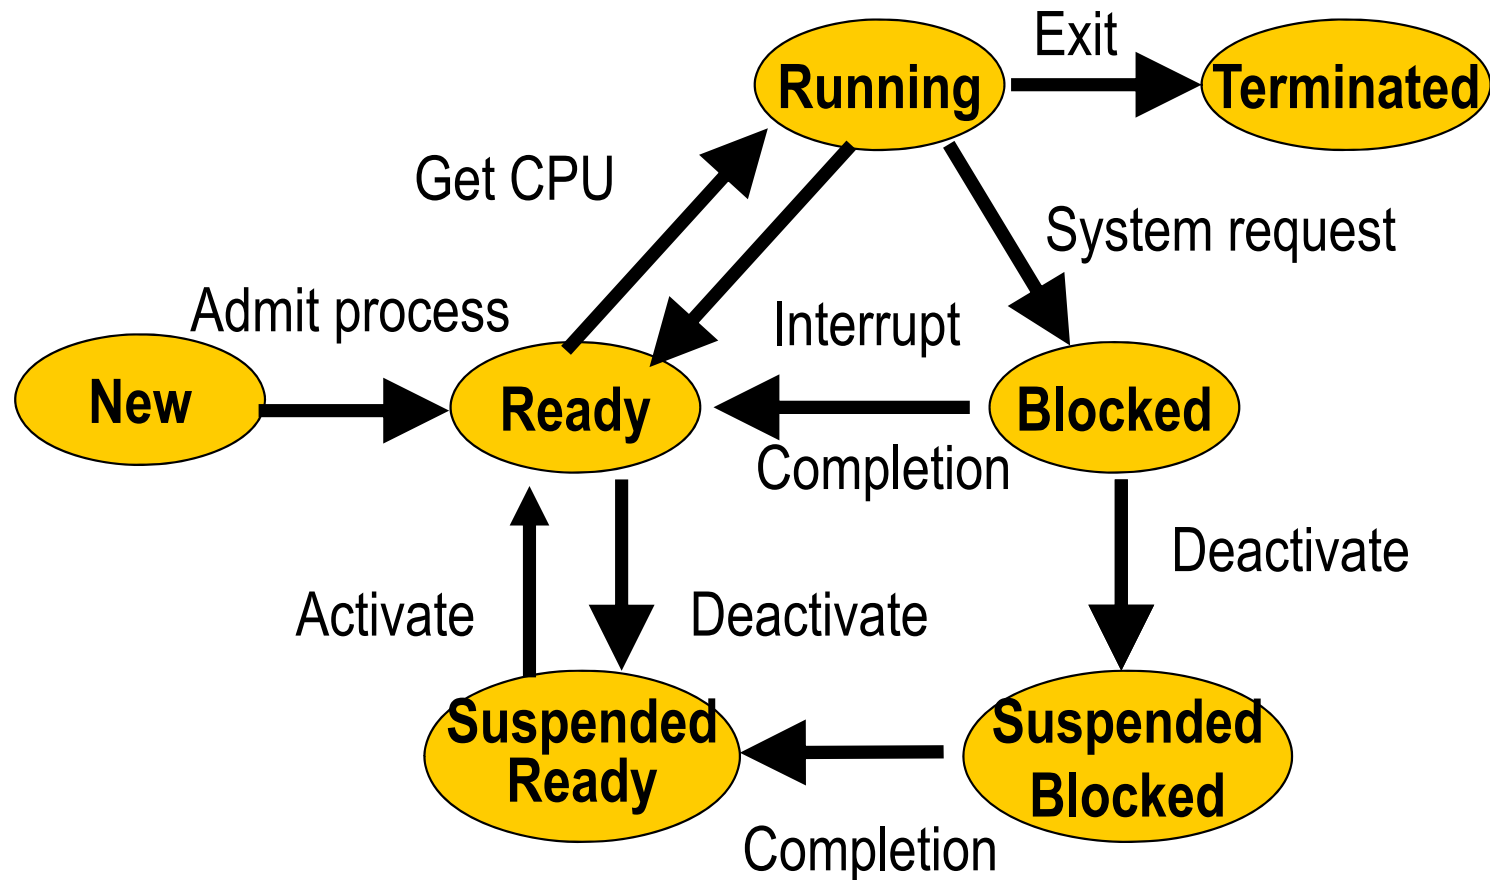

- System-wide table containing
 - ***Process identification*** (pid), *argument vector*, ...
 - ***Process current state***
 - ***Process priority and other CPU scheduling information***
 - A ***pointer*** to the remaining information.

Swapping




- Whenever the system is very loaded, we might want to expel from main memory or ***swap out***
 - Low priority processes
 - Processes that have been waiting for a long time for an external event
 - *User is out of the office*
- These processes are said to be ***swapped out*** or ***suspended***.

How it works




Suspended processes



- Suspended processes
 - Do not reside in main memory
 - Continue to be included in the process table
- Can distinguish between two types of suspended processes:
 - Waiting for the completion of some request (***blocked_suspended***)
 - Ready to run (***ready_suspended***).



A warning



- A system should ***not*** swap out ready processes unless their priority is ***very low***
- Otherwise swapping out ready processes can only be a ***desperate measure***