




COSC 3360/6310


Monday, February 22

Announcements




- First quiz now completely graded
 - Very slow process
 - Resolved to make next quiz self-grading
 - Answers to questions now posted on Prulu.com
- One week extension for the first assignment
- Second quiz will stay on Monday, March 1
- Do not be afraid to ask me assignment questions

Three honest reasons why you should ask instructor's help



- The programs will be graded by the TAs
 - Will never know how much help you get from the instructor
- You show that you have been working on the assignment
 - Make extensions more likely
- You will not be labelled for life as a bad programmer
 - We only remember final outcomes

Unfinished business



- Some people have not yet accepted the invitations I sent to their UH Mail accounts in January
- Need to reschedule a makeup exam.
 - Still time to catch up if you were unable to take the first quiz along with the other students.




Chapter II

Processes

Jehan-François Pâris
jfparis@uh.edu

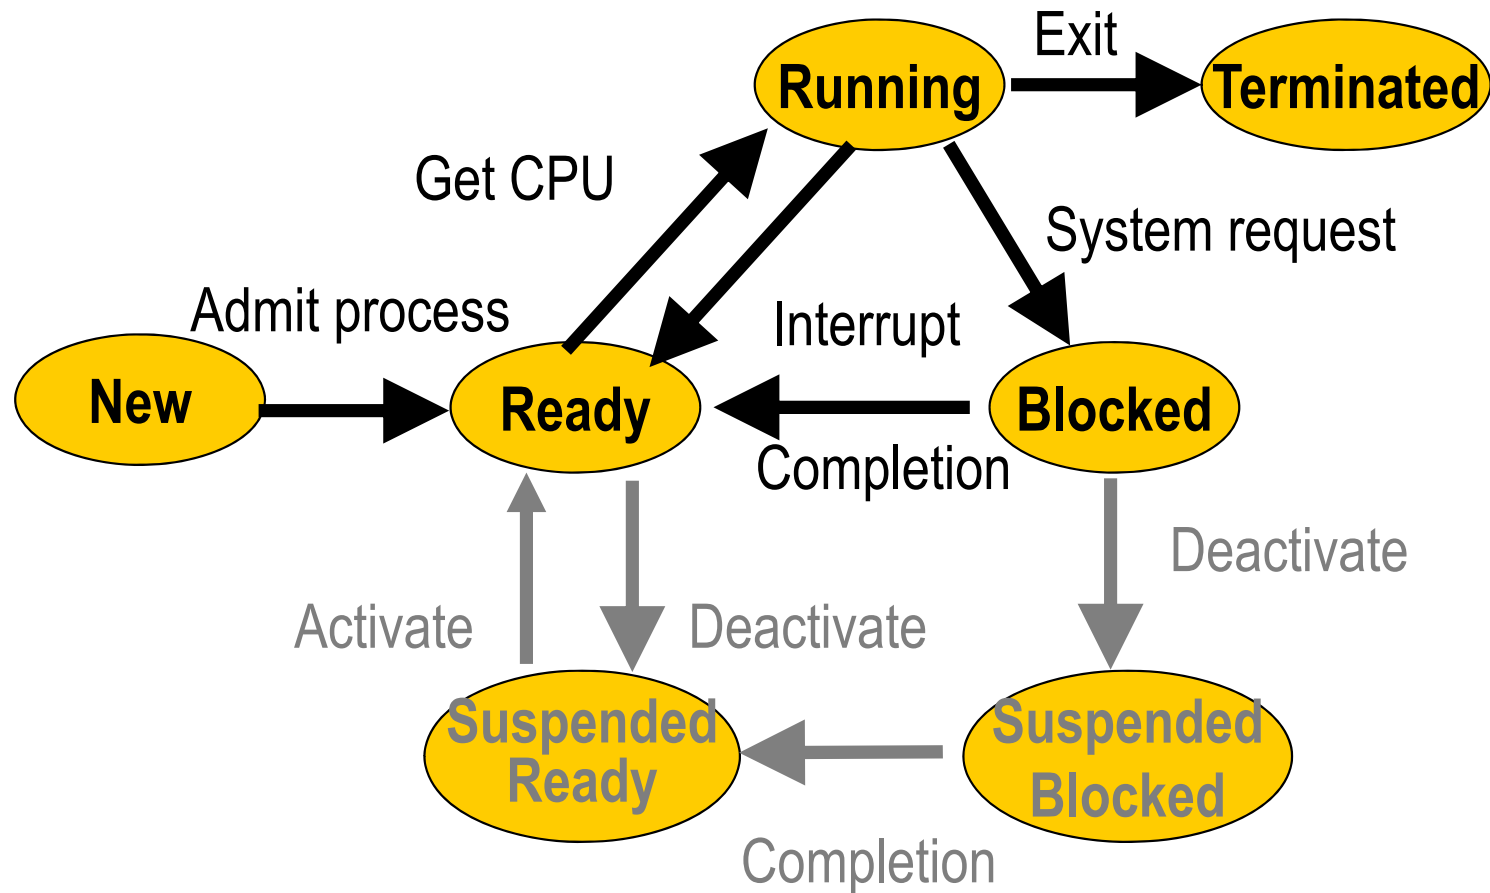




Chapter Overview



- Processes
- States of a process
- Operations on processes
 - **fork()**, **exec()**, **kill()**, **signal()**
- Threads and lightweight processes

How it works






Operations on processes

Process creation, deletion, ...




Operations on processes




- Process creation
 - `fork()`
 - `exec()`
 - The argument vector
- Process deletion
 - `kill()`
 - `signal()`

Process creation




- Two basic system calls
 - **fork()** creates a carbon-copy of calling process sharing its opened files
 - **execv()** overwrites the contents of the process address space with the contents of an executable file

fork() (I)




- First process of a system is created when the system is booted
- All other processes are forked by another process
 - Their ***parent process***
 - Said to be ***children*** of that process

fork() (II)

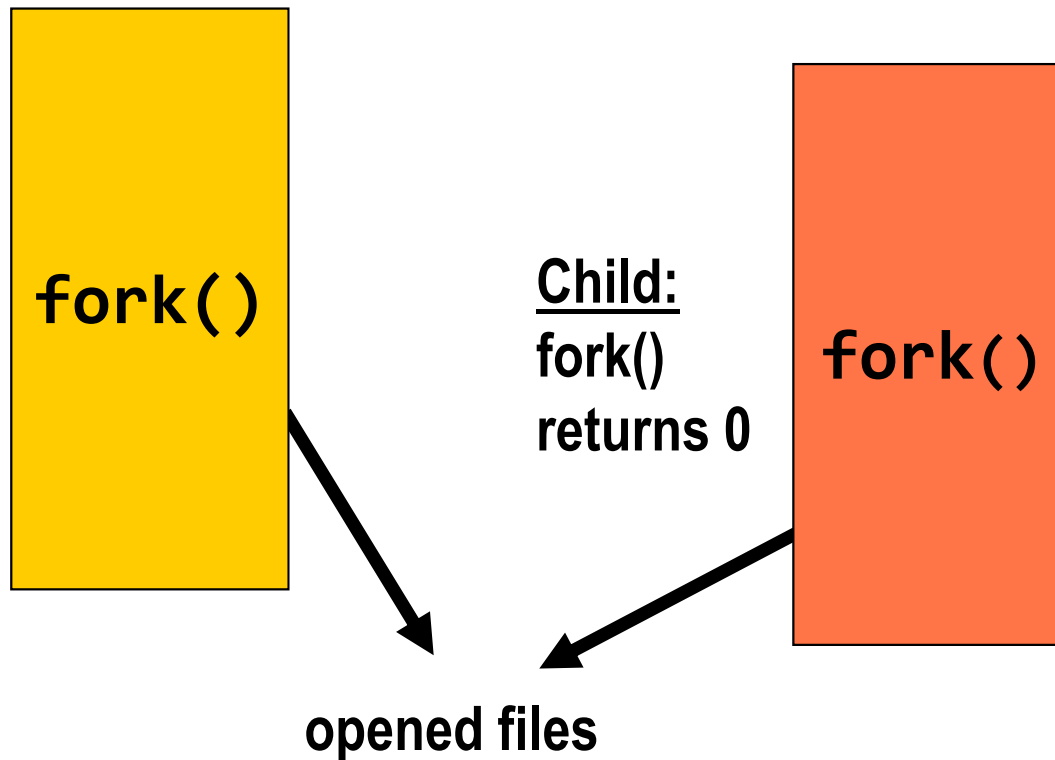


- When a process forks, OS creates an ***identical copy*** of forking process with
 - A ***new address space***
 - A ***new PCB***
- The ***only*** resources shared by the parent and the child process are the ***opened files***

fork() (III)




Parent:
fork()
returns
PID of
child



Child:
fork()
returns 0



First example




```
■ #include <iostream>
   using namespace std;
   main() {
       fork();
       cout << "Hello" << endl;
   } // main
```

will print two lines as **cout** will be executed by ***both*** the parent and the child

How it works





Second example




```
main() {  
    fork();  
    fork();  
    cout << "Hello" << endl;  
} // main
```

will print four lines as **cout** will be executed by the parent, its two children and its grandchild

How it works




Something smarter




```
int pid;
pid = fork();
if (pid == 0) {
    // child process
    ...
} else {
    // parent process
    ...
}
```

First simplification




```
int pid;
pid = fork();
if (pid == 0) {
    // child process
    ...
    _exit(0); // normal exit
} // if
// parent process continues
...
```

Second simplification



```
int pid;
if ((pid = fork())== 0) {
    // child process
    ...
    _exit(0); // normal exit
} // if
// parent process continues
...
```


Waiting for child completion



- **wait(0)**
 - Waits for the completion of any child
 - No wait if any child has already completed
- **while (wait(0) != kidpid)**
 - Waits for the completion of a specific child identified by its ***pid***




An example (I)




- ```
#include <iostream>
#include <sys/types.h>
#include <sys/wait.h>
using namespace std;
```

## An example (II)



```
■ main() {
 int pid;
 if((pid = fork()) == 0) {
 cout << "Hello !" << endl;
 _exit(0);
 } // child
 wait(0);
 cout << "Goodbye!" << endl;
} // main
```


# Notes



- Unix keeps in its process table all processes that have terminated but their parents have not yet waited for their termination
  - They are called ***zombie processes***
- The statement  
`while (kidpid != wait(0));`  
is a loop with an ***empty body***




# Putting everything together



```
int kidpid;
if ((kidpid = fork())== 0) {
 // child process
 ...
 _exit(0); // normal exit
} // if
// parent waits for child
while (wait(0) != kidpid);
...
```


Must use the while loop if the process has already forked other children

# exec




- Whole set of `exec()` system calls
- Most interesting are
  - `execv(pathname, argv)`
  - `execve(pathname, argv, envp)`
  - `execvp(filename, argv)`
- All `exec()` calls perform the same two tasks
  - Erase current address space of process
  - Load specified executable

# execv




- `execv(pathname, argv)`
  - `char pathname[]`
    - ***full pathname*** of file to be loaded:  
`/bin/ls` instead of `ls`
  - `char argv[][]`
    - the ***argument vector***:  
passed to the program to be loaded

# Argument vector (I)




- An array of pointers to the individual argument strings
  - `arg_vector[0]` contains the name of the program *as it appears in the command line*
  - Other entries are parameters
  - End of the array is indicated by a **NULL** pointer


# Argument vector (II)



- `char argv[][];`
- `char **argv;`




# execve() and execvp()




- **execve(pathname, argv, envp)**
  - Third argument points to a list of environment variables
- **execvp(argv[0], argv)**
  - Lets user specify a command name instead of a full pathname
  - Looks for **argv[0]** in list of directories specified in environment variable **PATH**

# Putting everything together



```
int pid;
if ((pid = fork()) == 0) {
 // child process
 ...
 execvp(filename, argv);
 _exit(1); // exec failed
} // if
while (pid != wait(0));
// parent waits
...
```


# Observations (I)



- Not cheap
  - `fork()` makes a ***complete copy*** of parent address space
    - ***Very costly*** in a virtual memory system
  - `exec()` thrashes that address space
- Best solution is copy-on-write (COW)




# Copy-on-write



Parent and child share same address space




When either of them modifies a page,  
other gets its ***own copy*** of original page



COW of original page


# Copy-on-write as a lazy approach



- Copy-on-write postpones address space copying until it is actually needed
  - Do the strict minimum
- ***Lazy approach***
  - Betting that very little copying will be actually needed
    - An `execv()` will quickly follow
- Opposite is ***eager approach***




# Observations (II)




- Neither `fork()` nor `exec()` affect opened file descriptors
  - They remain unchanged
- Important for Unix I/O redirection mechanism

# How this happened




- Fork was not that expensive on a minicomputer with a 16-bit address space
  - Never had to copy more than **64KB**
- Using a fork/exec allowed a very easy implementation of I/O redirection
  - After the **fork()** thus in the child
  - Before the **exec()** while parent is still in control

# A very basic shell (I)




```
for (;;) {
 parse_input_line(argv);
 if built_in(argv[0]) {
 do_it(arg_vector);
 continue;
 } //built_in command
 path = find_path(argv[0]);
```

## A very basic shell (II)




```
if ((pid = fork()) == 0) {
 // put here I/O
 // redirection code
 execv(path, argv);
 _exit(1); // execv failed
} //child process
if (interactive())
 while (wait(0) != pid);
} // main for loop
```

# Comments



- Shell built-in commands include
  - **exit**  
terminates the shell
  - **cd**  
changes current directory
- Commands are assumed to be interactive
  - ***Non-interactive*** commands end with an “&”

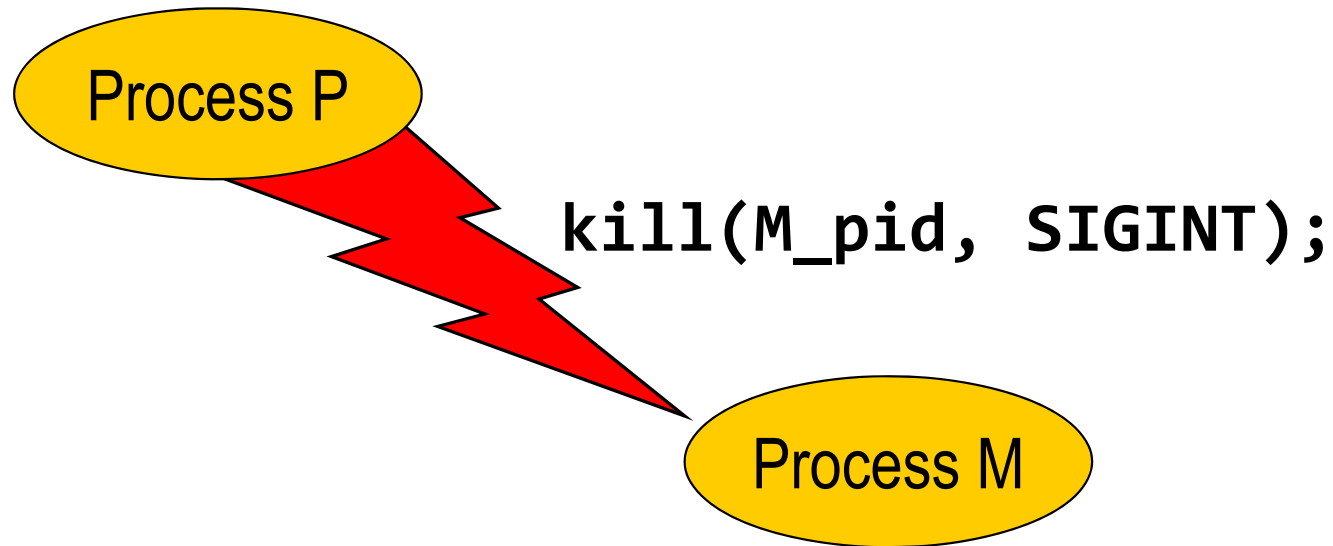

# Terminating a process (I)



- Sending a signal:
  - `kill()` has two arguments
    - The ***process id*** of the receiving process
    - A ***signal name*** or a ***signal number***
- `#include <signal.h>`  
`kill(this_pid, this_signal);`
- Process receiving the signal will ***terminate***




# Terminating a process (II)





What should I do? **AARGH!**

# Catching a signal (I)



- The process receiving signal can ***catch*** it by using **signal()**
  - Will not terminate
- **signal(a\_signal, catch\_it);**
  - **catch\_it** function is executed when **a\_signal** signal is received.
- The ninth signal, **SIGKIL**, cannot be caught.

## Catching a signal (II)



Process is now **shielded** by `signal()` call