




COSC 3360/6310


Wednesday, February 24

Announcements



- First assignment is now due on March 3 anywhere on earth
- Second quiz will be on Monday, March 1
 - Will cover chapter 2 of class notes
 - Three video lectures
 - F. Wednesday February 10
 - G. Monday February 22
 - H. Wednesday February 24

More announcements



- Some people have not yet accepted the invitations I sent to their UH Mail accounts in January
 - Please check your UH mail account (jsbach@uh.edu)
- Videos of last two lectures are now online
 - Had problems with MS Teams default permissions




Chapter II

Processes

Jehan-François Pâris
jfparis@uh.edu

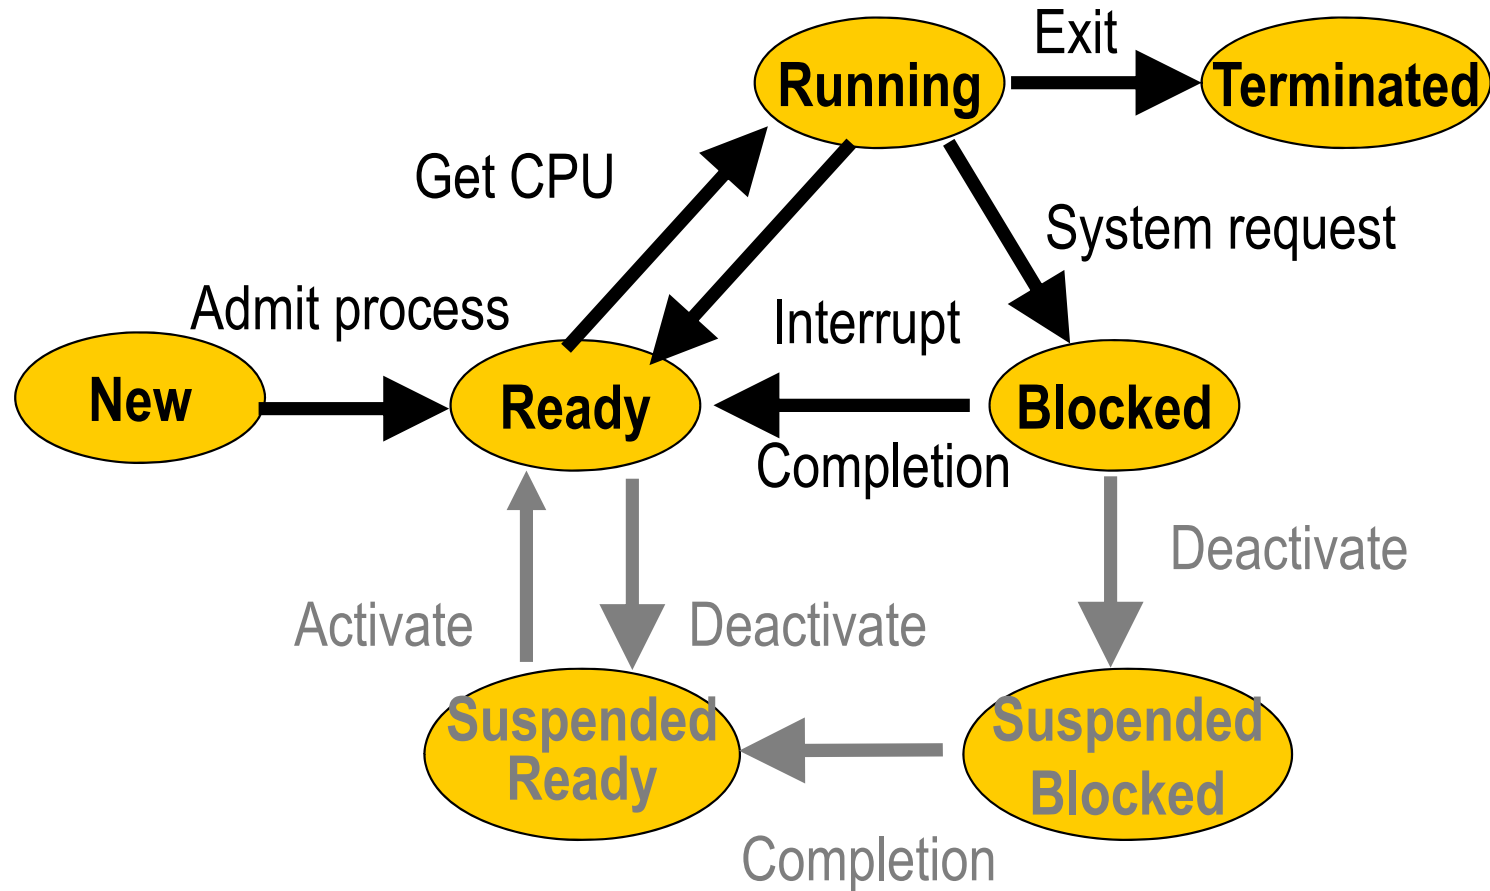




Chapter Overview



- Processes
- States of a process
- Operations on processes
 - **fork()**, **exec()**, **kill()**, **signal()**
- Threads and lightweight processes

How it works






Operations on processes

Process creation, deletion, ...




Operations on processes



- Process creation
 - `fork()`
 - `exec()`
 - The argument vector
- Process deletion
 - `kill()`
 - `signal()`


What will this program print out?



```
■ #include <iostream>
using namespace std;
main() {
    if ((pid = fork()) == 0) {
        //child
        cout << "Hello" << endl;
    }
    // parent
    cout << "Goodbye" << endl;
} // main
```




Answer



- Hello
Goodbye
Hello

or


- Goodbye
Hello
Goodbye



Lightweight processes/threads

Kernel supported threads, user-level threads, POSIX threads (pthreads)

Limitations of processes




- Single threaded server:
 - Processes one request at a time

```
for (;;) {  
    receive(&client, request);  
    process_request(...);  
    send(client, reply);  
} // for
```




A basic question



- ***What does a server do when it does not process client requests?***




Three good answers



- Nothing
- It waits for client requests
- It “sleeps”
 - ***Blocked state is sometimes called the sleep state***


The problem



- Most client requests involve disk accesses
 - File servers
 - Authentications servers
- When this happens, the server remains in the BLOCKED state
 - Cannot handle other customers' requests
- Could end doing nothing most of the time
- ***Poor throughput***




An analogy




- *In most fast-food restaurants, counter employees process customer orders one order at a time.*
- *Not be possible in a traditional restaurant*
 - *A server that would only be able to wait on one table at a time would be idle most of the time.*

A first solution



```
int pid;
for (;;) {
    receive(&client, request);
    if ((pid = fork()) == 0) {
        process_request(...);
        send(client, reply);
        _exit(0); // done
    } // if
} // for
```

The good and the bad news



- ***The good news:***


- Server can now handle several user requests in parallel

- ***The bad news:***

- **fork()** is a ***very expensive*** system call
 - Has to create a new address space




A better solution




- Provide a faster mechanism for creating cheaper processes:
 - ***Lightweight processes***
 - ***Threads***

How?



- Lightweight processes and threads ***share the address space of their parent***
 - ***No need to create a new address space***
 - Most expensive step of `fork()` system call

Is it not dangerous?




- ***To some extent because***

- No memory protection inside an address space
- Lightweight processes can now interfere with each other

- ***But***



- All lightweight process code is written by the same team

General Concept (I)




- A ***thread*** or ***lightweight process***
 - Does ***not*** have its ***own address space***
 - Shares it with its parent and other peer threads in the same address space (***task***)
- Each thread has a ***program counter***, a ***set of registers*** and its ***own stack***.
 - *Everything else is shared*

General Concept (II)




- A regular process (single-threaded)




- A process containing several threads

Implementation




- Threads and LWPs can either be
 - ***Kernel supported:***
 - Mach, Linux, Windows NT and after
 - ***User-level:***
 - Original pthread library, ...

Kernel-Supported Threads (I)




- Managed by the kernel through system calls
- One process table entry per thread
- This is the best solution for *multiprocessor architectures*
 - Kernel can allocate ***several processors*** to a ***single multithreaded task***

Kernel-Supported Threads (II)



- Supported by Mach, Linux, Windows NT and more recent systems
- ***Performance Issue:***
 - Switching between two threads in the same task involves a system call
 - Results in ***two context switches***

Linux Threads




- **clone(fn, stack, flags)**

where


- **fn** specifies function to be executed by new thread or process
- **stack** points to the stack it will use
- **flags** is a set of flags specifying various options
 - **CLONE_VM** for threads
 - Regular process if **CLONE_VM** is missing

User-Level Threads (I)



- User-level threads are managed by procedures **within** the task address space
 - The *thread library*
- One process table entry per task/address space
 - Kernel is not even aware that process is multithreaded

User-Level Threads (II)



- Can be retrofitted into an OS lacking thread support
 - Portable thread libraries
- ***No performance penalty:***
 - Switching between two threads of the same task is done cheaply within the task
 - Same cost as a procedure call



User-Level Threads (III)



- ***Programming issue:***

- Each time a thread does a ***blocking system call***, kernel will move the ***whole process*** to the ***blocked state***
 - It does not know better
- Must then use ***non-blocking*** system calls
 - *Complicates programmer's task*

User-Level Threads (IV)




`sleep (5) ;`

Kernel

**Process wants to sleep for 5 seconds:
Should be moved it to the blocked state**

POSIX Threads




- POSIX threads, or ***pthread***s, started as pure user-level threads managed by the POSIX thread library
 - Gained later ***some kernel support***
- Ported to various Unix and Windows systems (***Pthreads-win32***).
- Function names start with **pthread_**
- Calls tend to have a complex syntax

An Example (I)



FYI




```
#include <pthread.h>  
static int count[2];
```

Static variables are shared by all threads
Other variables are stored on the private stack of each thread.

An Example (II)




FYI



```
void *child(void *arg) {  
    int index;  
    index = (int) arg;  // required  
    for(;;) {  
        printf("Child count: %d\n",  
               ++count[index]);  
        sleep(1); // one second delay  
    } // for loop  
} // child
```

An Example (III)




```
int main() {  
    thread_t tid; // thread id  
    int i = 0;  
    pthread_create(&tid, NULL,  
                  child, (void *) i);  
    // pthread will execute  
    // "child" function
```

**NULL stack address specifies
a new stack "anywhere"**

An Example (IV)




FYI



```
i++; // now i == 1
while (count[i] < 12) {
    printf("Parent count: %d\n", ++count[i]);
    sleep(1); // one second delay
} // while loop
return 0;
} // main
```


Understanding pthread_create()



- **pthread_create()** has four arguments
 - **&tid**
 - Placeholder for **thread_id**
 - **NULL**
 - Stack address of new stack
 - **NULL** means “anywhere”
 - **start_function**
 - Void pointer to a function
 - **(void *) arg**
 - Sole argument passed to **start_function**




Comparing the approaches




<i>Feature</i>	<i>Kernel threads</i>	<i>User-level threads</i>
<i>Portability</i>		
<i>Multiprocessing</i>		
<i>Performance</i>		
<i>Ease of use</i>		

Comparing the approaches




<i>Feature</i>	<i>Kernel threads</i>	<i>User-level threads</i>
<i>Portability</i>		☑
<i>Multiprocessing</i>		
<i>Overhead</i>		
<i>Ease of use</i>		

Comparing the approaches




<i>Feature</i>	<i>Kernel threads</i>	<i>User-level threads</i>
<i>Portability</i>		☑
<i>Multiprocessing</i>	☑	
<i>Overhead</i>		
<i>Ease of use</i>		

Comparing the approaches




<i>Feature</i>	<i>Kernel threads</i>	<i>User-level threads</i>
<i>Portability</i>		☑
<i>Multiprocessing</i>	☑	
<i>Overhead</i>		☑
<i>Ease of use</i>		

Comparing the approaches



<i>Feature</i>	<i>Kernel threads</i>	<i>User-level threads</i>
<i>Portability</i>		☑
<i>Multiprocessing</i>	☑	
<i>Overhead</i>		☑
<i>Ease of use</i>	☑	

Conclusion



- No clear winner between kernel-supported and user-level threads
- Solaris (from Sun, now taken over by Oracle)
 - Supports both ***user-level threads*** and ***kernel threads***
 - Lets programmers combine them as they need